

INTRODUCTION TO TRANSACTION CONCURRENCY

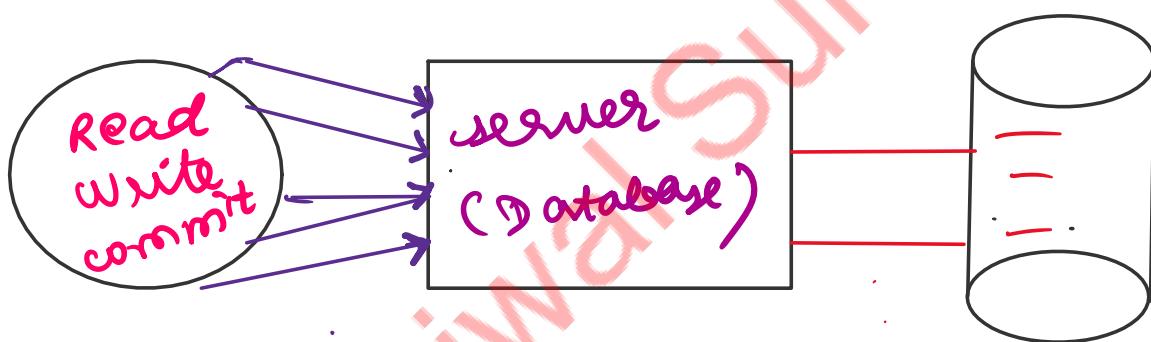
15 October 2023 12:03

Transactions

↳ It is a set of operations used to perform a logical unit of work.

A transaction generally represents a change in database.

work → withdraw money



Trans for $A \rightarrow B$

database Access
change

Example : $A = 1000$, $B = 2000$

Transfer 500 from $A \rightarrow B$.

① $R(A) \rightarrow A = 1000$

④ $R(B) \rightarrow B = 2000$

② $A = A - 500$

⑤ $B = B + 500$

③ $W(A) \rightarrow A = 500$

⑥ $W(B) \rightarrow B = 2500$

... and 1

③

WCA) \leftarrow $\pi = \emptyset$

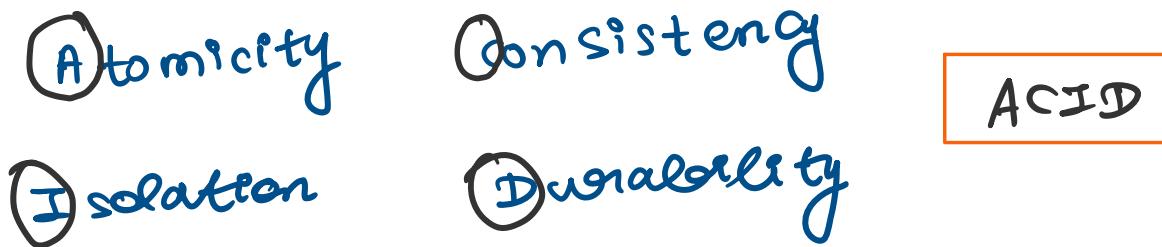
[All in RAM only]

Finally ⑦ commit \leftarrow save it in the database

[R \rightarrow read, W \rightarrow write]

ACID PROPERTIES

15 October 2023 13:27



Atomcity

either all or none

In a transaction, if any

one step fails be fall commit

ROLLBACK is done.

→ as if no statement in the transactions were executed.

A failed transaction cannot be recovered, it will always restart.

consistency

→ before the transaction starts, and after the transaction ends, the sum of the total money should be the same.

10 000

be the same.

Eg : $A = 2000$
 $B = 3000$



$R(A) \rightarrow A = 2000$

$A = A - 1000$

$w(A) \rightarrow A = 1000$

$R(B) \rightarrow B = 3000$

$B = B + 1000$

$w(B) \rightarrow B = 4000$

commit

$A \xrightarrow{1000} B$

Before transaction:

sum = $A + B$

= $2000 + 3000$
= 5000

After transaction:

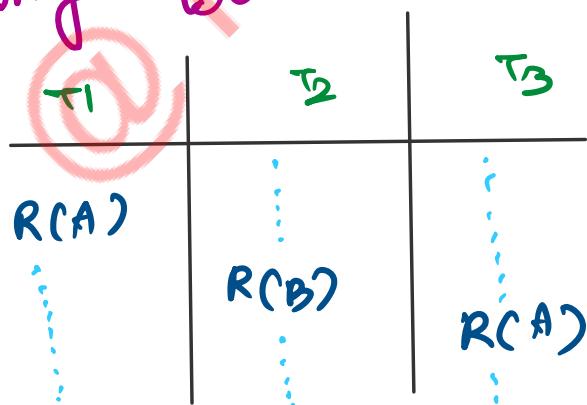
sum = $1000 + 4000$

= 5000

(same)

Isolation

many transactions running parallelly



$T_1 \rightarrow T_2$ (or)

$T_2 \rightarrow T_1$

can a parallel schedule be converted
into a serial schedule in real time
(P. No. conversion)?

into a transaction
only conceptually?

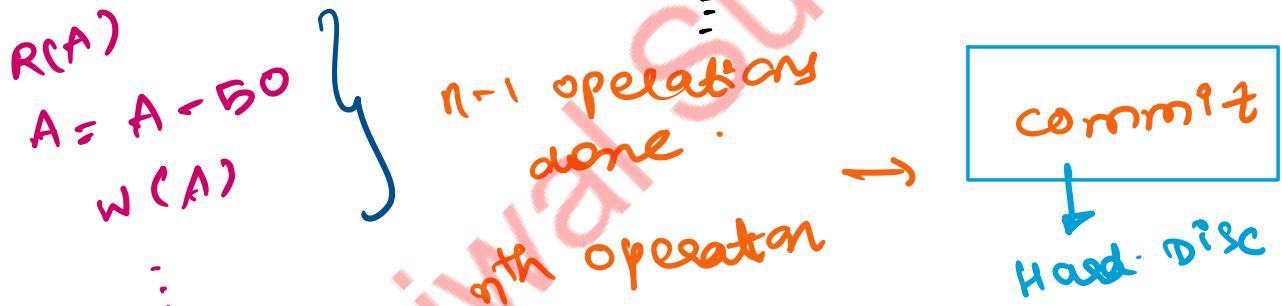
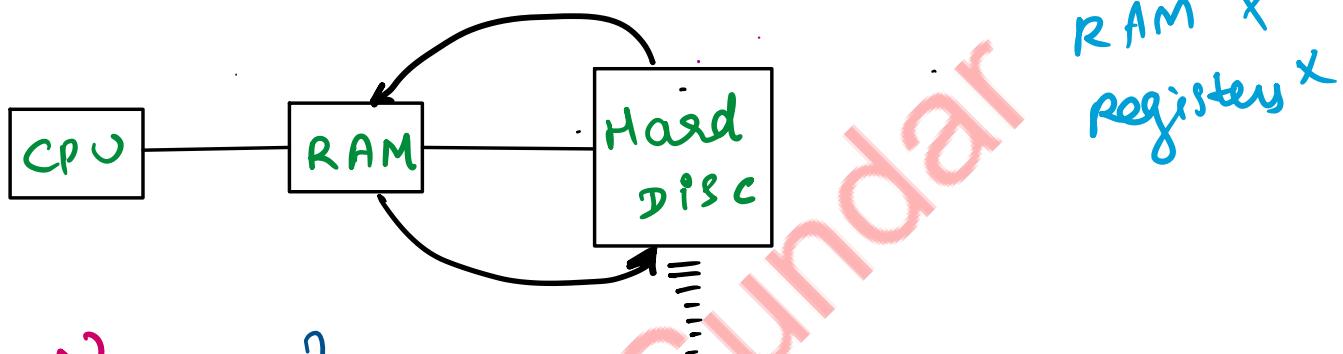
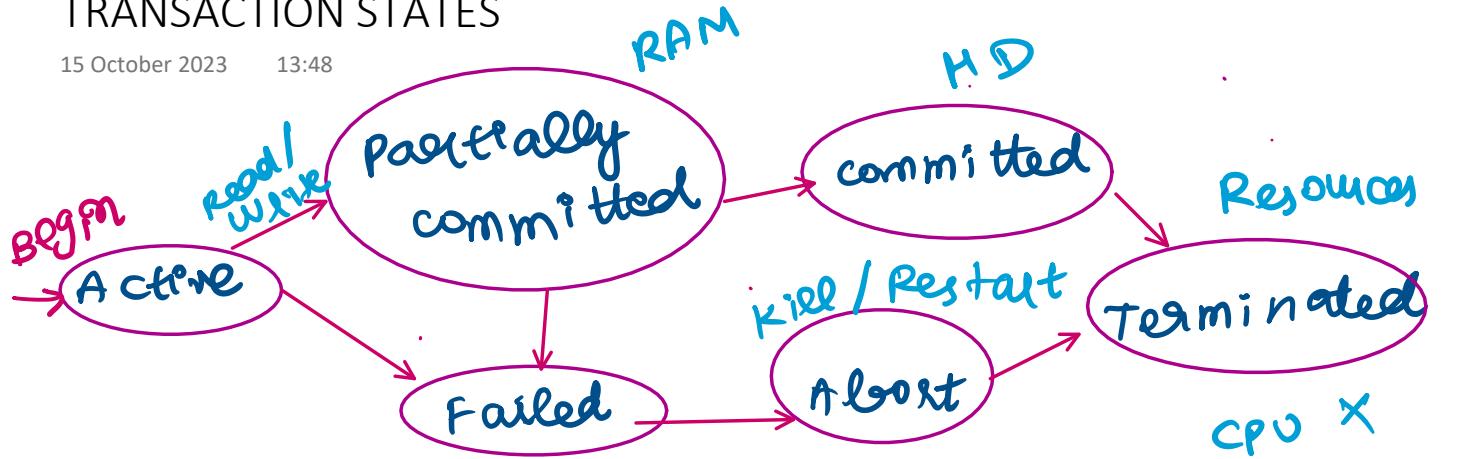
[serial schedule \rightarrow always consistent]

Durability

↓
whatever changes are performed on
the database must be permanent.

TRANSACTION STATES

15 October 2023 13:48



SCHEDULE

15 October 2023 16:22

Schedule

↳ It is the chronological execution sequence of multiple transactions.

$T_1 \quad T_2 \quad T_3 \quad \dots \quad T_n$

serial schedule :

→ (eg ATM)

$T_1 \quad T_2 \quad T_3$

Any transaction that starts, completes fully without any interruptions. → so secure.

$\xrightarrow{T_1} \xrightarrow{T_2} \xrightarrow{T_3}$

serial transactions are always consistent.

But the problem is waiting time → even though T_1, T_2, T_3 arrived at the same time → T_2 & T_3 had to wait for T_1 , then T_3 had to wait for T_2 .

parallel schedule :

$T_1 \quad T_2 \quad T_3$
1

↓
eg online banking system

eg online ~~order~~^{order} system

Here waiting time ↓

Performance ↑

very high

throughput ↑

no. of transactions
executed / time

@ Prajwal Sunkari

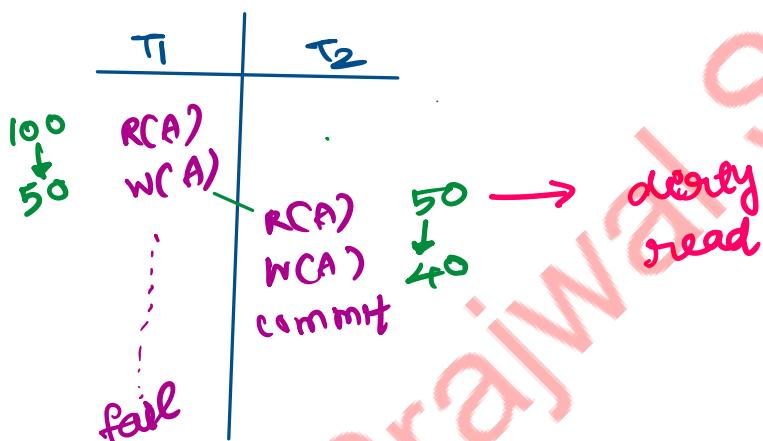
PROBLEMS IN CONCURRENCY

15 October 2023 16:33

- * dirty read
- * incorrect summary
- * lost update
- * unrepeatable read
- * phantom read

(parallel schedule)

Dirty Read / Uncommitted Read / RAW



Incorrect summary

T ₁	T ₂	$A = 1000, B = 1000, C = 1000$
$\text{read}(A)$ $A = 1000$	$s = 0$ $\text{avg} = 0$ $\text{read}(C)$ $s = s + C$	$s = 0 \text{ (sum)}$ $\text{avg} = 0$ $T_2 R : C = 1000$ $s = 1000$ $T_1 R : A = 1000$

new(m)

$A = A - 50$

write(A)

read(A)

$S = S + A$

read(B)

$S = S + B$

$avg = S/3$

commit

read(B)

$B = B + 50$

write(B)

commit

T₁ W : $A = 950$

T₂ R : $A = 950$

$S = 1950$

T₂ R : $B = 1000$

$S = 2950$

$avg = 983.33$

T₁ R : $B = 1000$

T₂ W : $B = 1050$

Lost Update

Alice

begin

select

Q: 7
L: 5

Bob

begin

Q
↓
quantity

L
↓
likes

select

update qty = 6

update qty = 10

lost \leftarrow commit : $Q = 6$
 $L = 5$

commit : $Q = 10$

lost \leftarrow

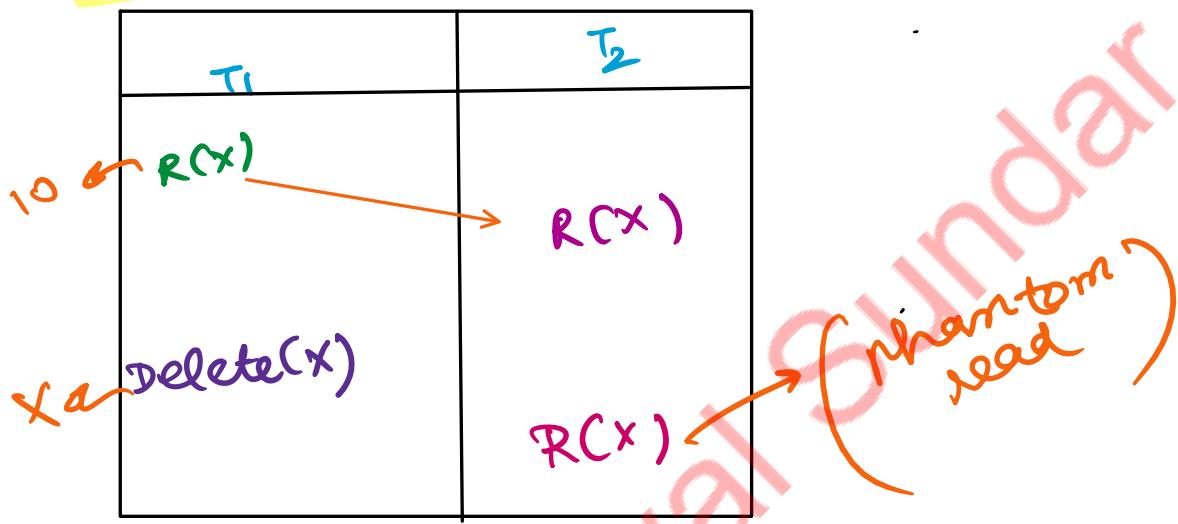
$$L = 5$$

commit : $Q = 10$

$$L = 5$$

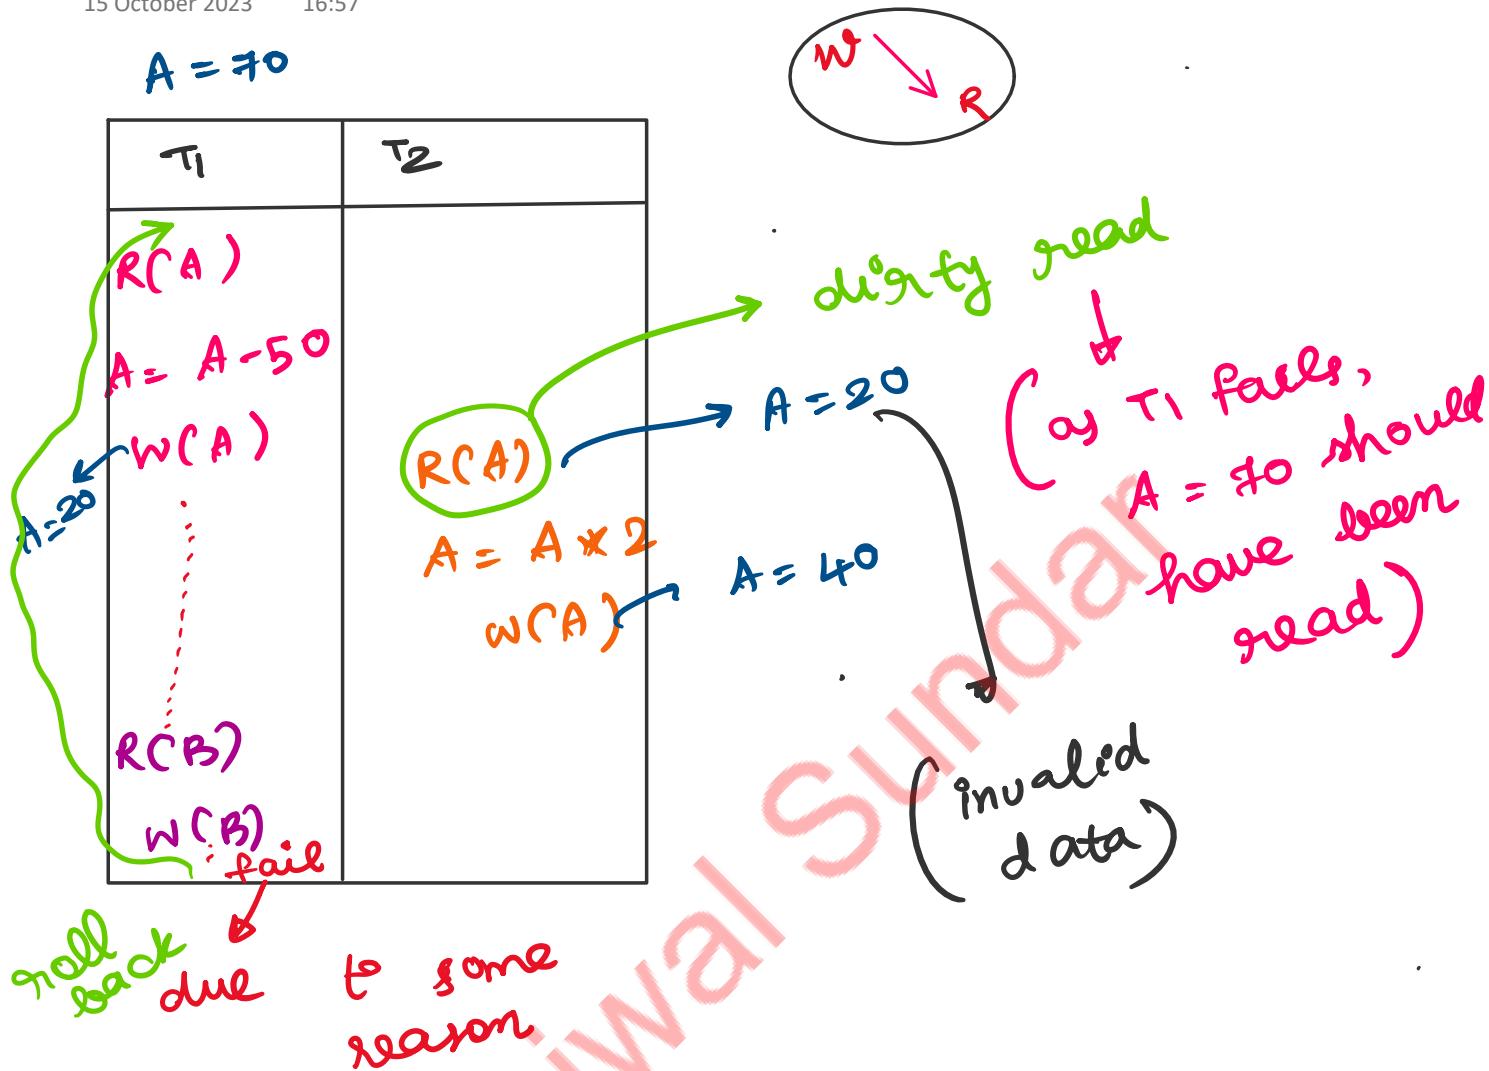
↓
overwrites
previous commit

Unrepeated read



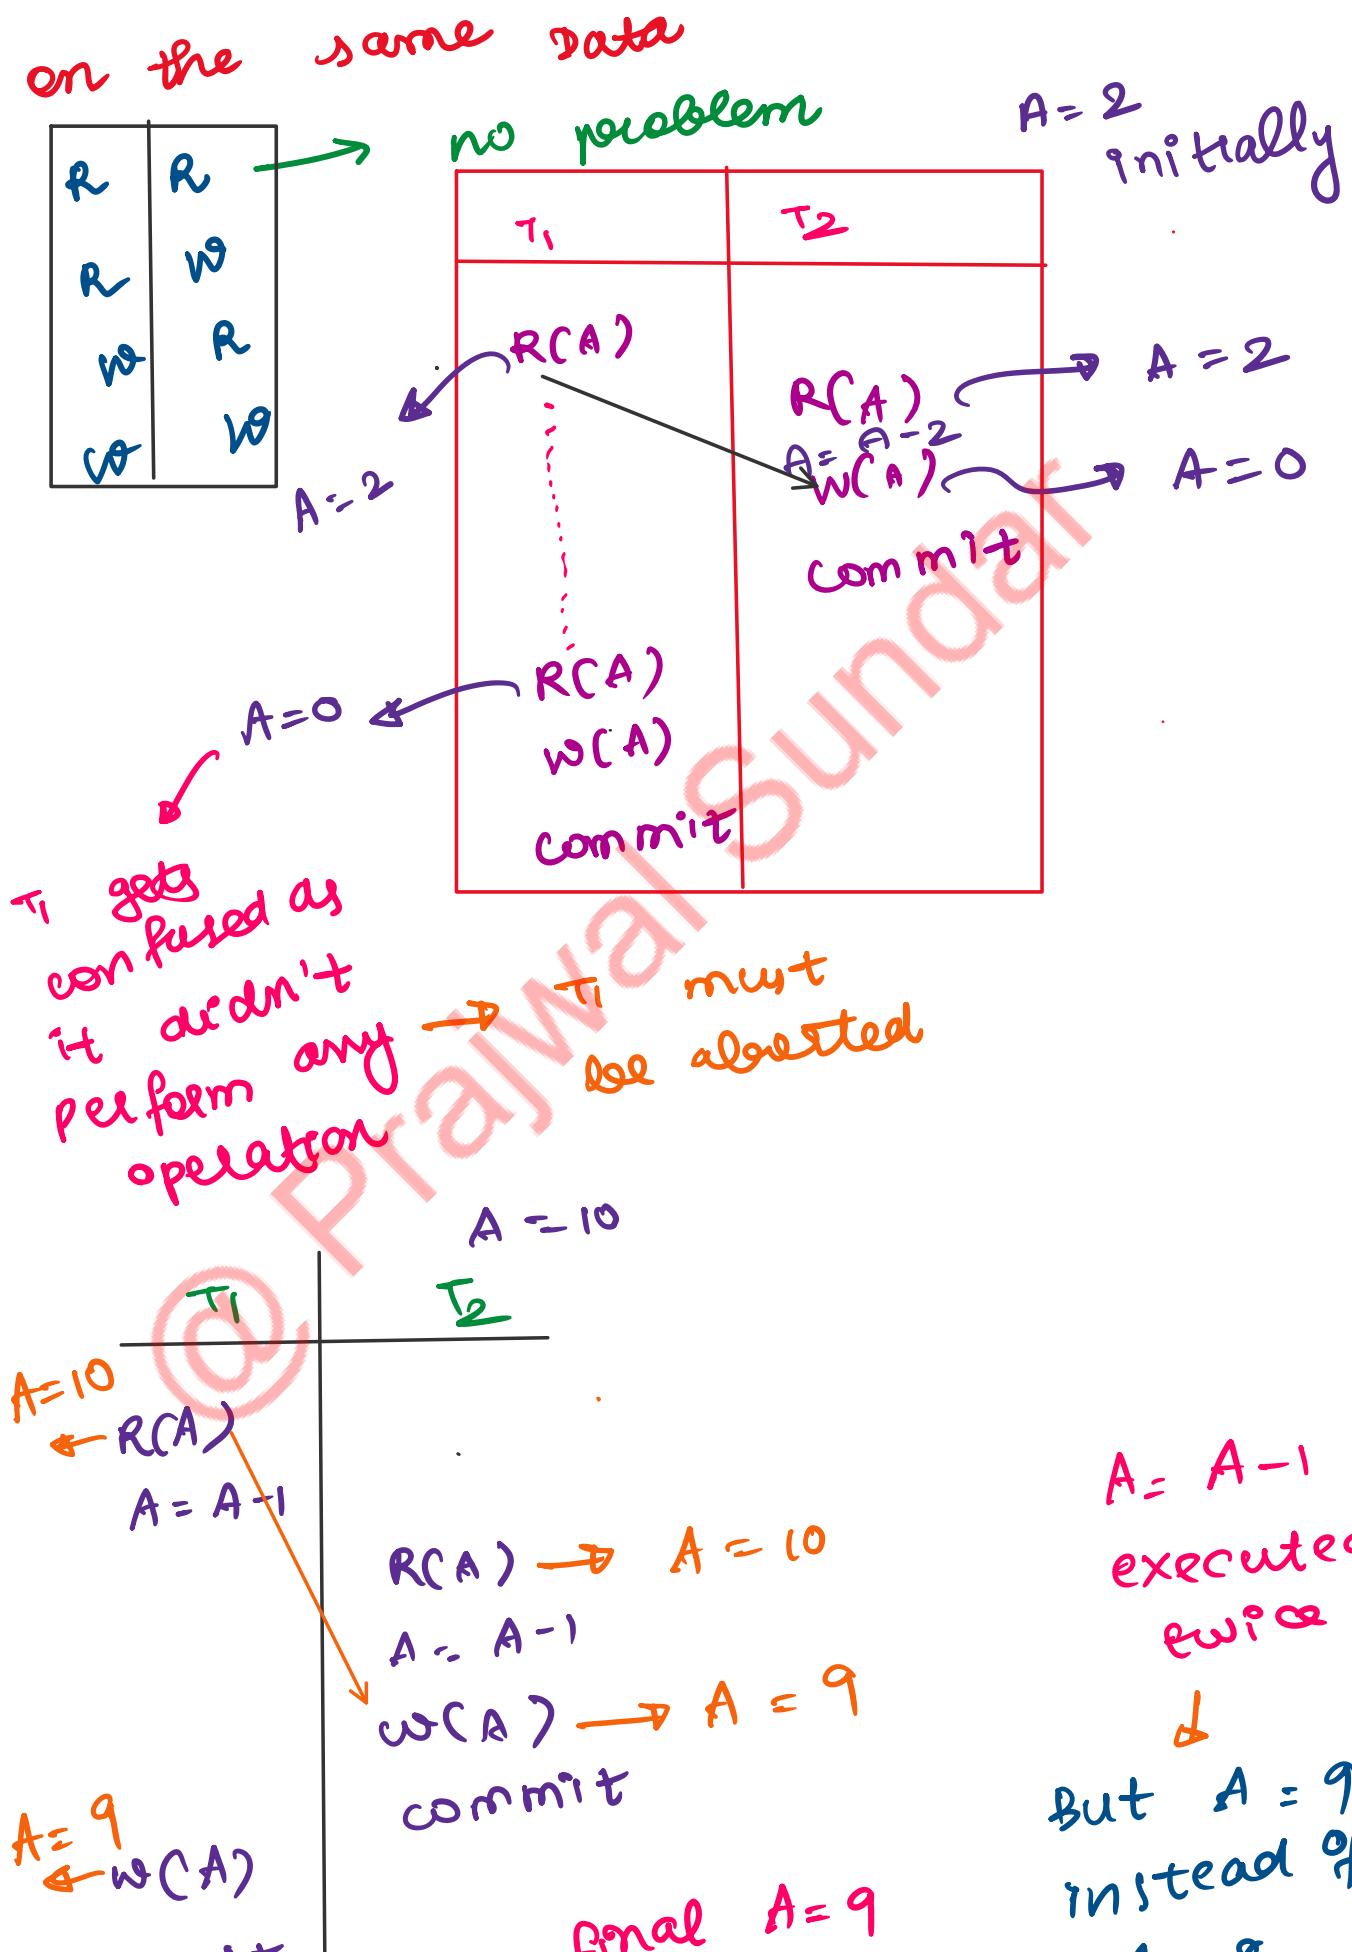
WRITE-READ CONFLICT [DIRTY READ PROBLEM]

15 October 2023 16:57



READ-WRITE CONFLICT [UNREPEATABLE READ]

15 October 2023 17:04



$w(A)$
commit

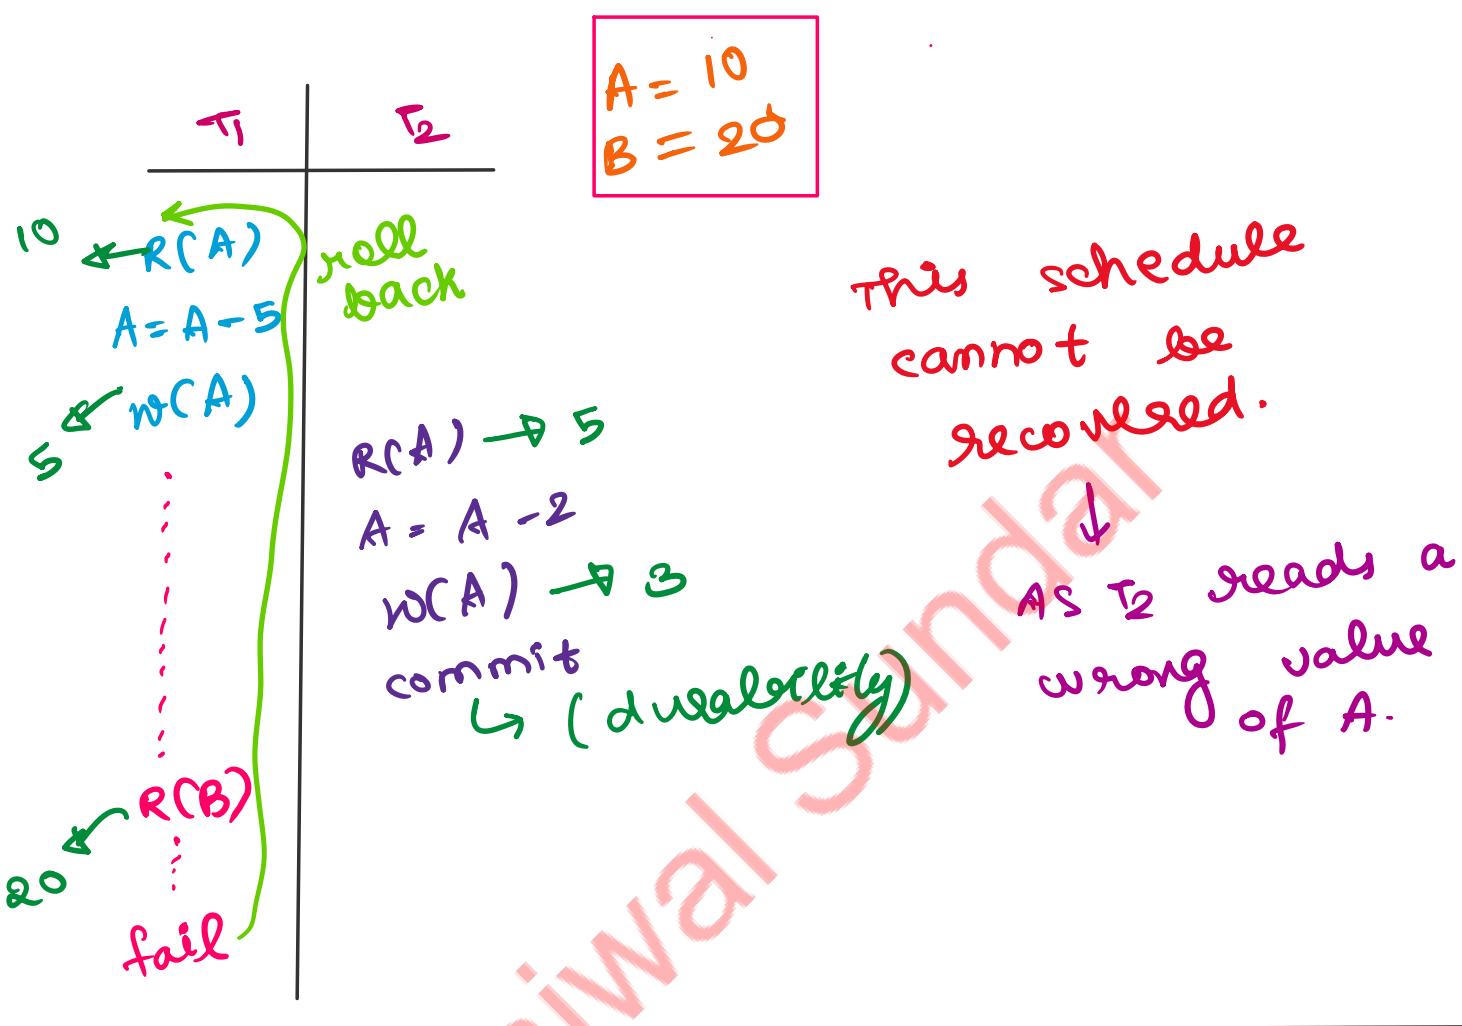
final $A = 9$

instead of
 $A = 8$

@ Prajwal Sundar

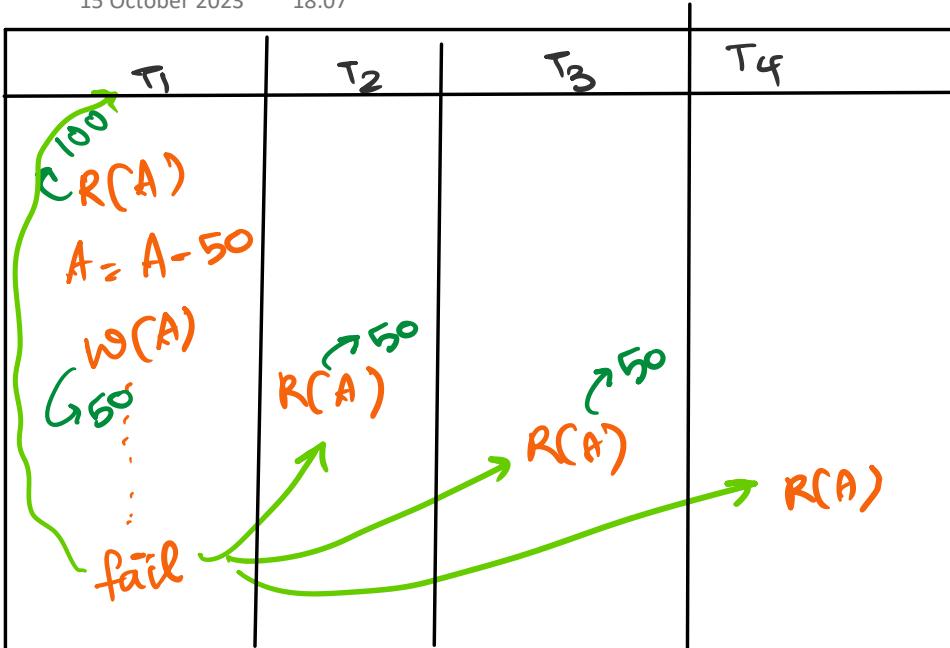
IRRECOVERABLE SCHEDULE \rightarrow (re recoverability)

15 October 2023 18:02



CASCADING VS CASCADELESS SCHEDULE

15 October 2023 18:07



while performing roll back
→ rollback all connected transactions
[cascade]

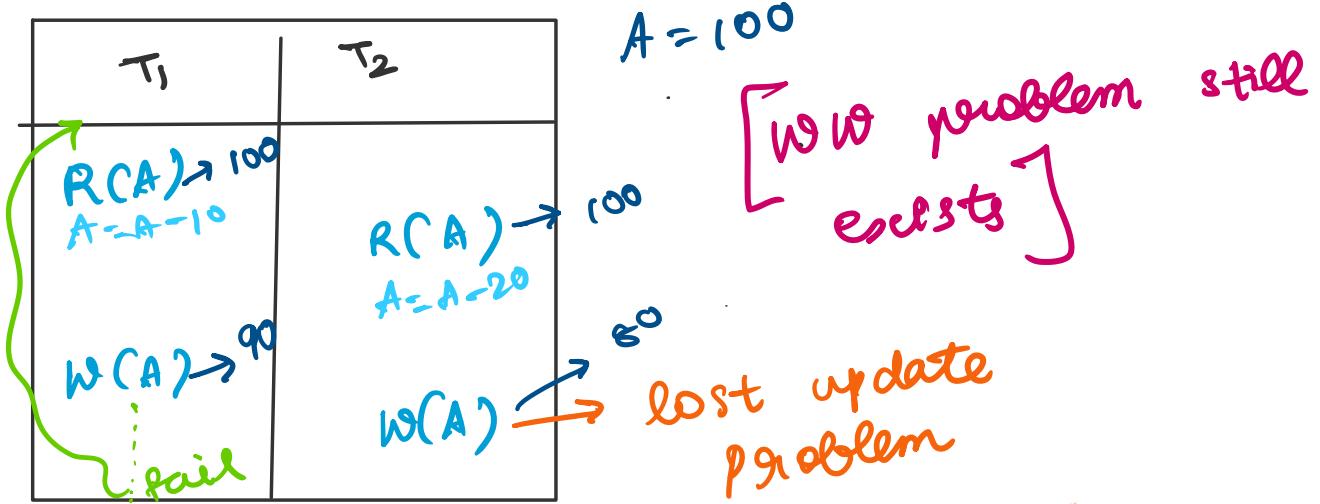
↳ CPU utilization ↓
performance is degraded

T_1	T_2	T_3
$R(A)$ $w(A)$	$R(A) \times$	$R(A) \times$

T_1	T_2	T_3
$R(A)$ $w(A)$ commit	$R(A)$	$R(A)$

while T_1 is operating on data A, do not allow other transactions to access it until either:
* T_1 commits (or) * T_1 fails.

$$A = 100$$



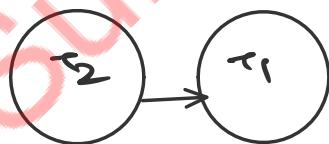
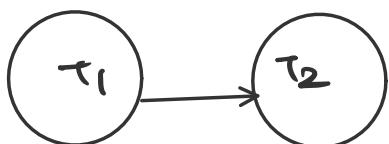
SERIALIZABILITY

15 October 2023 18:30

↳ ability to make a schedule serializable

T_1	T_2
$R(A)$	
$w(A)$	$R(A)$

T_1	T_2
	$R(A)$
	$w(A)$



These 2 schedules are already serial schedules \rightarrow no need to serialize / check serializability.

T_1	T_2
$R(A)$	
	$R(A)$

↓
(parallel module)

checking for serializability

↓
check if an equivalent schedule of the given parallel schedule can be obtained or not -

Types \rightarrow conflict
 \rightarrow new

(parallel schedule)

types → view

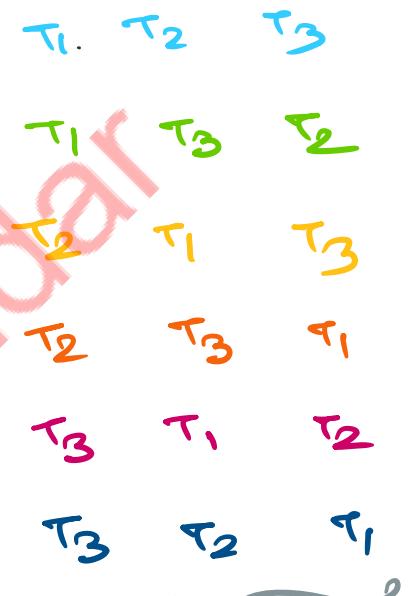
[convert to $\tau_1 \rightarrow \tau_2$ (OR) $\tau_2 \rightarrow \tau_1$]

Another example :

[$3 = 6$ cases]

τ_1	τ_2	τ_3
$R(B)$ $w(B)$	$R(A)$ $w(A)$	$R(A)$ $w(A)$
	$w(B)$	

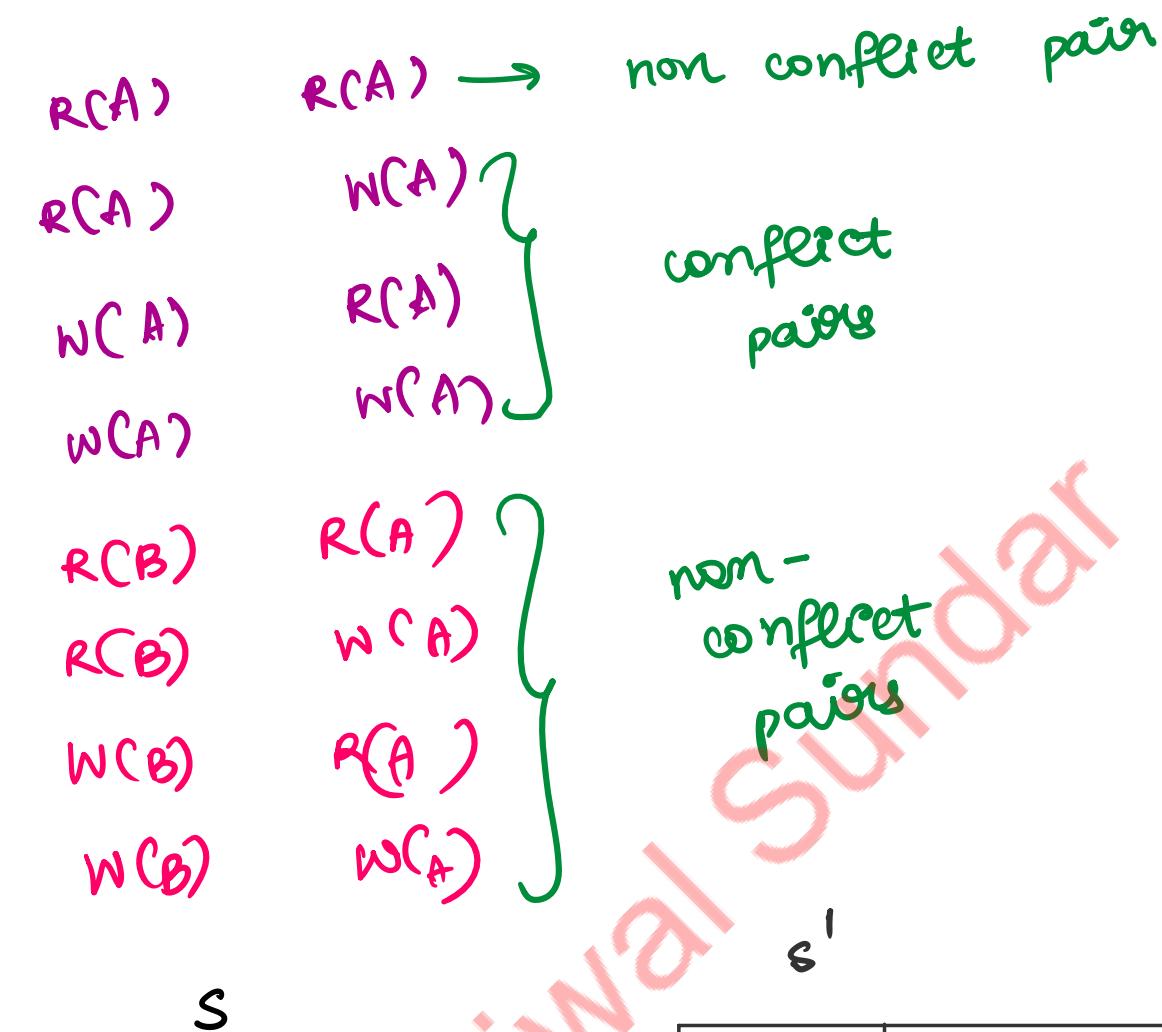
parallel schedule



if the given parallel schedule can be converted into any one of these @ combinations, it is serializable.

CONFLICT EQUIVALENT SCHEDULE

15 October 2023 18:56



T_1	T_2
$R(A)$	
$W(A)$	
$R(A)$	
$W(A)$	
$R(B)$	

Method = swap adjacent non-conflict pairs

$s \rightarrow s'$

IS
 $s \equiv s' ?$
 ↴
 check

T_1	T_2
-------	-------

T_1 T_2

$R(A)$

$W(A)$

$R(B)$

$R(A)$

$W(A)$

↓

T_1	T_2
RCA)	
w(A)	
R(B)	RCA)
	w(A)

swap →

T_1	T_2
RCA)	
w(A)	
R(B)	
	R(A)
	w(A)

serial schedule
obtained.

$S \equiv S' \rightarrow$ obtained

$\therefore S$ is conflict serializable.

check for conflict serializability:

T_1	T_2
RCA)	
w(A)	
w(A)	RCA)
R(B)	

conflicting pairs
&
no change in positions
possible

→ not conflict serializable

CONFLICT SERIALIZABILITY

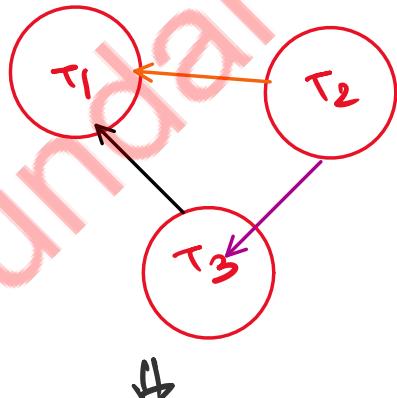
15 October 2023 19:29

check conflict pairs in other transactions and draw edge

[RW, WR, WW]

T_1	T_2	T_3
$R(x)$ $R(z)$ $w(x)$ $w(z)$	$R(y)$ $R(z)$ $w(z)$	$R(y)$ $R(x)$ $w(y)$

precedence graph :



check for loop / cycle in the graph

If → conflict serializable

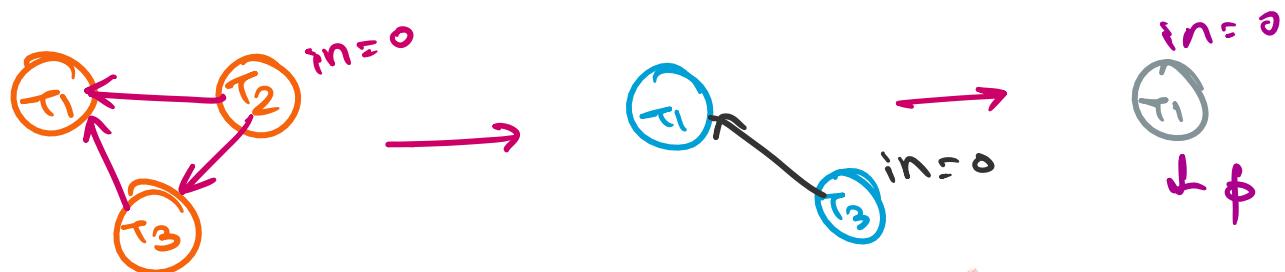
↓
serializable

↓
consistent

To find which order, use :
topological sort

remove vertex with in-degree = 0

[choose vertex with in-degree = 0
and erase it from the graph]



serial order : $T_2 \rightarrow T_3 \rightarrow T_1$

Equivalent serial schedule :

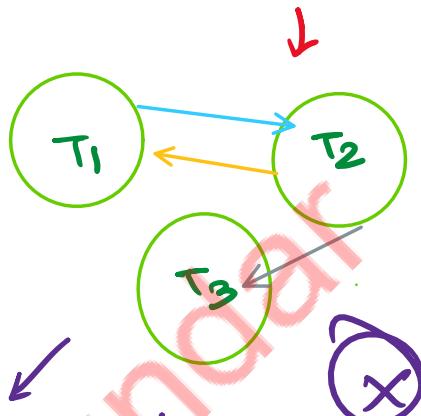
T_1	T_2	T_3
$R(x)$ $R(z)$ $w(x)$ $w(z)$	$RC(y)$ $RC(z)$ $w(z)$	$RC(y)$ $R(x)$ $w(y)$

WHY VIEW SERIALIZABILITY?

15 October 2023 19:49

T_1	T_2	T_3
$R(A)$	$N(A)$	$w(A)$

check whether given schedule is conflict serializable or not



To topological sort is not possible
loop detected.
It is non conflict serializable.
may / may not be serializable.

If no loop / cycle \rightarrow CS \rightarrow serial
 \rightarrow consistent

To check serializability
view serializability is used

$A = 100$		
T_1	T_2	T_3
$R(A)$	$A = A - 40$ $w(A)$	
$A = A - 40$ $w(A)$		$A = A - 20$ $w(A)$

T_1	T_2	T_3
$R(A)$ $A = A - 40$ $w(A)$	$A = A - 40$ $w(A)$	$A = A - 20$ $w(A)$



Final value of A is the same.

They look same \rightarrow view equivalent ✓
But not conflict equivalent ✗

@ Prajwal Sundar

SHARED-EXCLUSIVE LOCKING

15 October 2023 20:24

shared lock (S)

↳ If transaction locked data item in shared mode, then it is allowed to read only.

Exclusive lock (X)

↳ If transaction lock data item is in exclusive mode, then it is allowed to read and write both.

(T₁)

s(A)

r(A)

u(A)

(T₂)

x(A)

r(A)

w(A)

u(A)

↓
shared lock
@
exclusive lock

s request

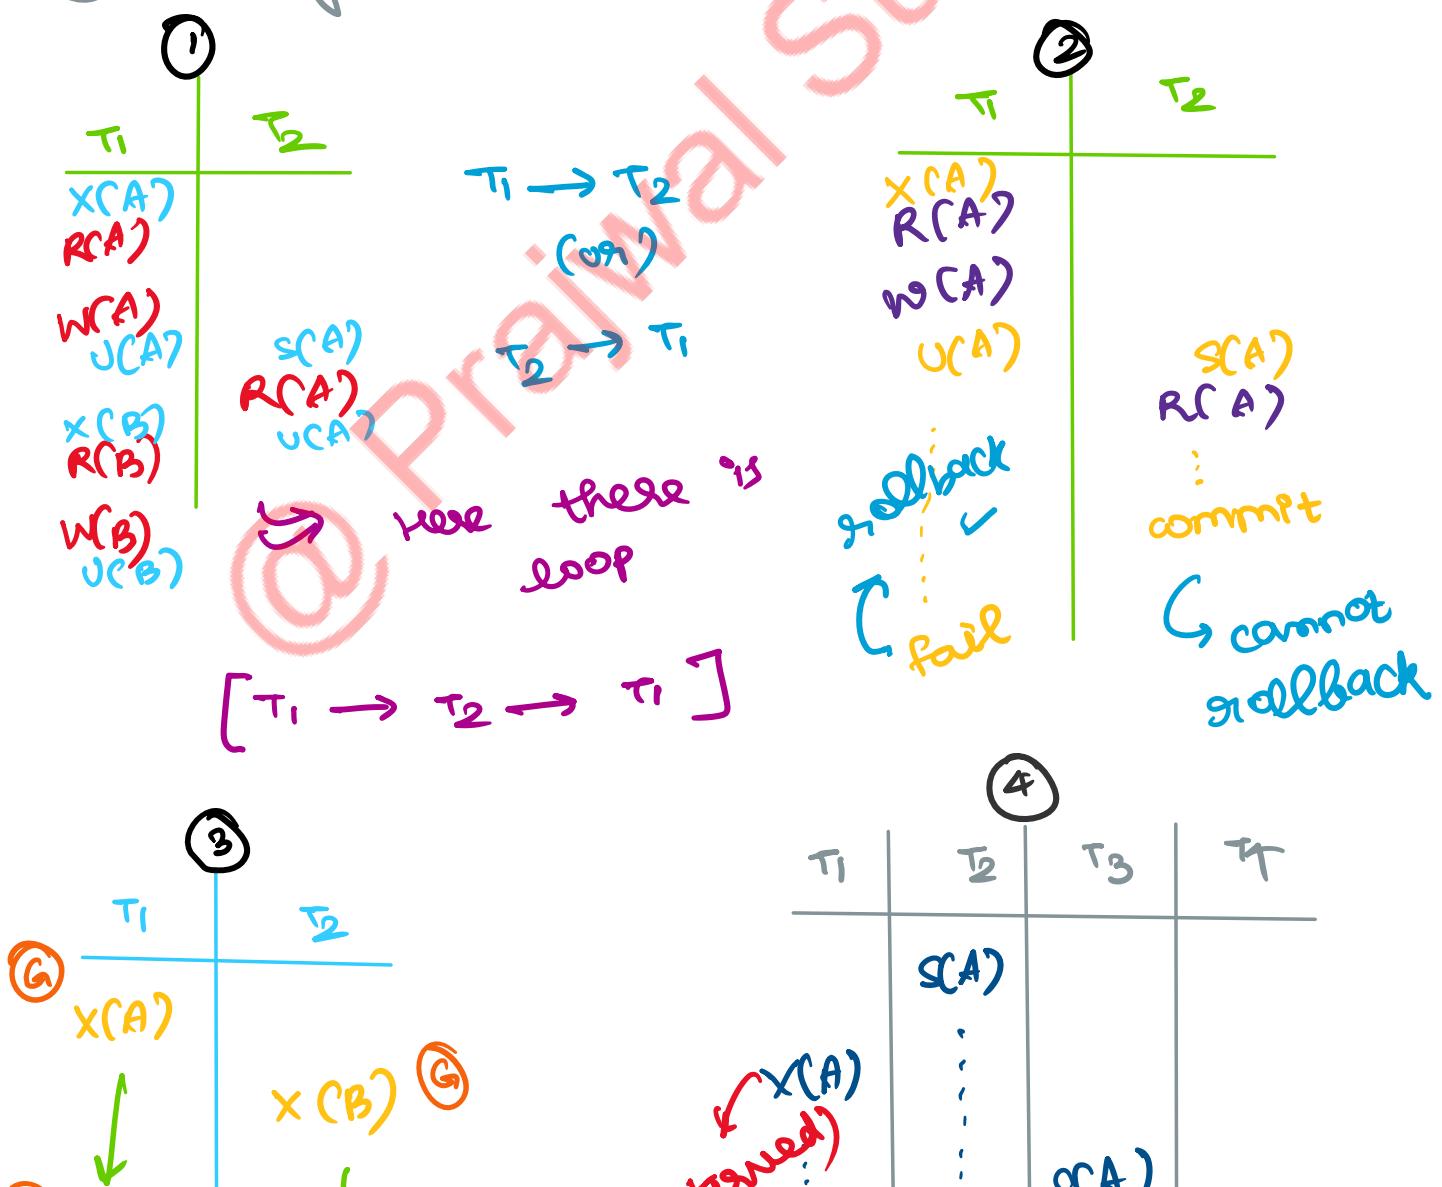
Grant S	s
Yes	No
No	No

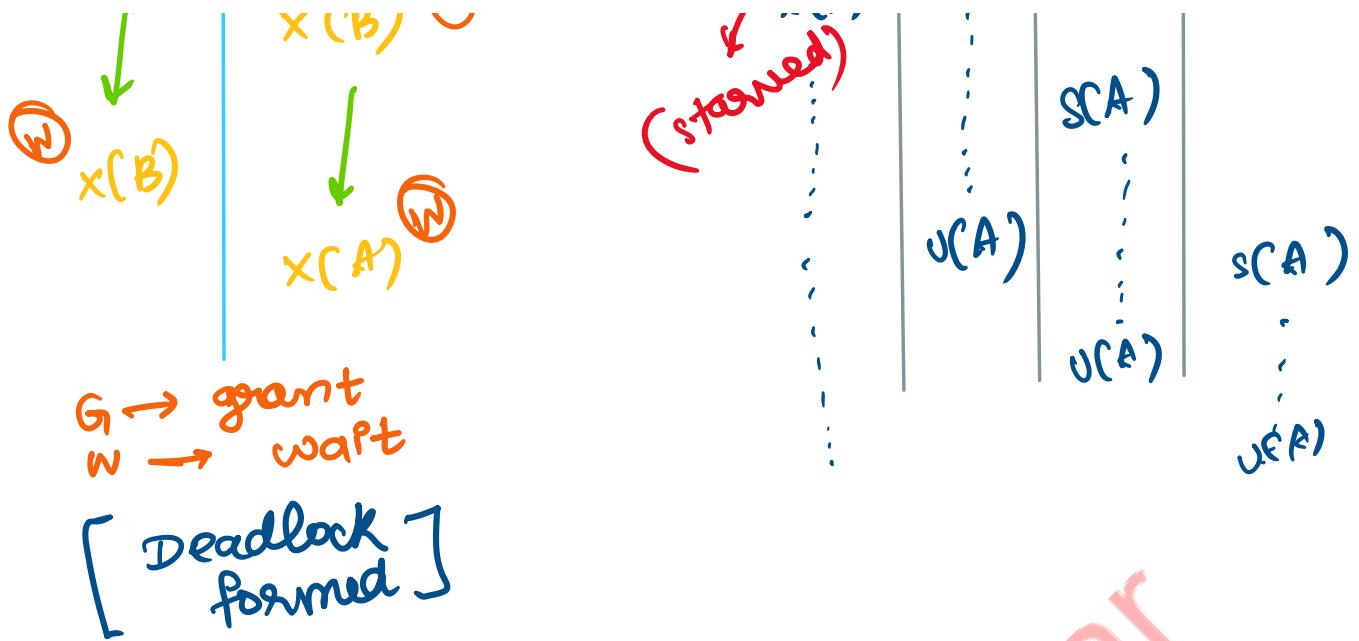
DRAWBACKS IN SHARED / EXCLUSIVE LOCKING PROTOCOL

04 November 2023 11:00

Problems in S/X Locking :

- ① may not be sufficient to produce only serializable schedule.
- ② may not be free from non-availability.
- ③ may not be free from deadlock.
- ④ may not be free from starvation.

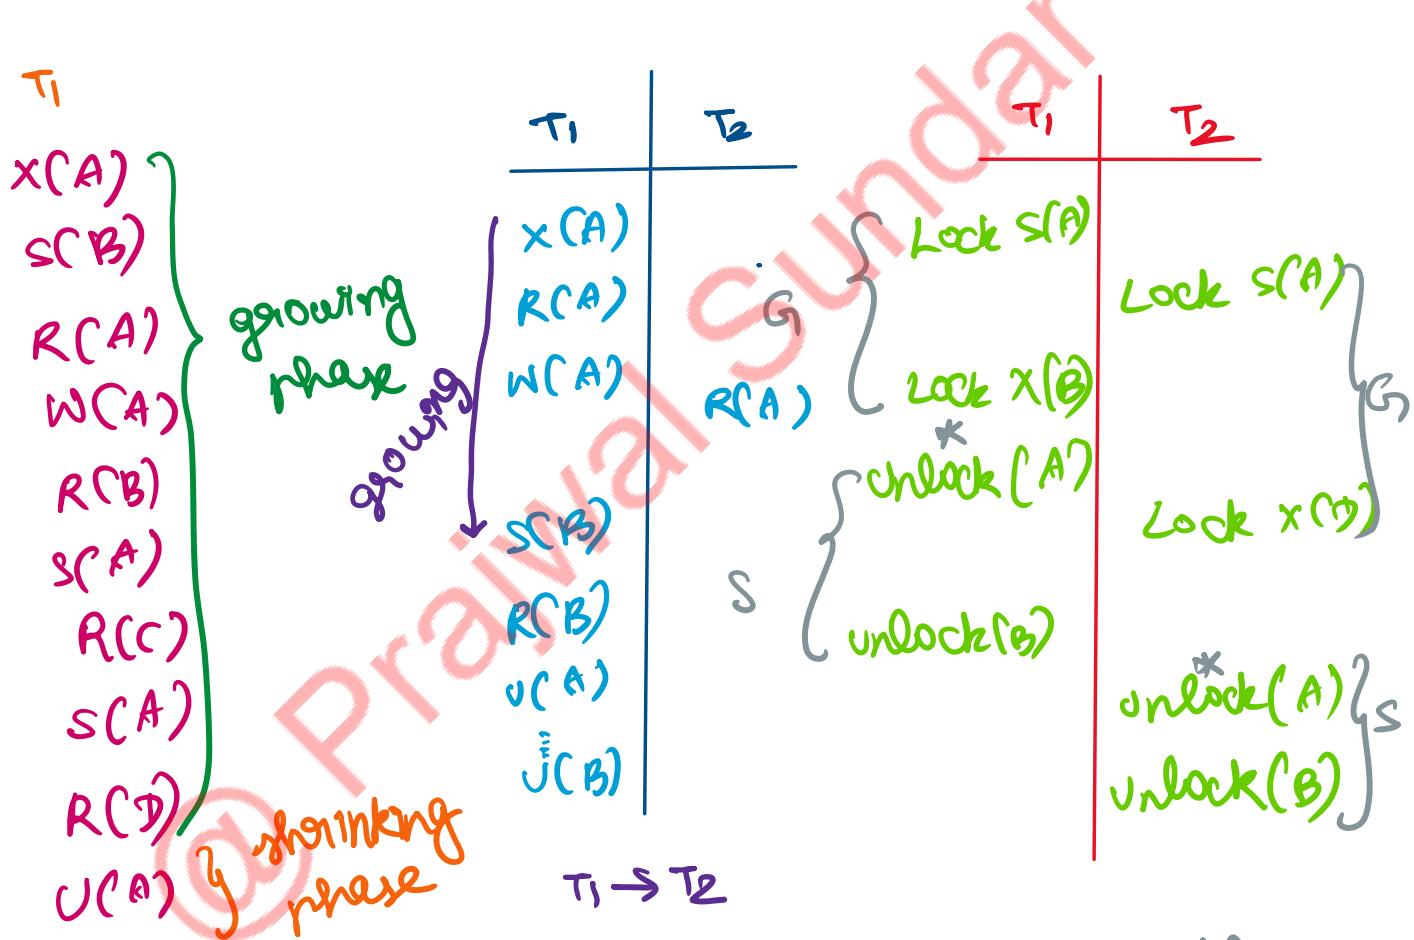




2 PHASE LOCKING

04 November 2023 11:14

- [Growing phase
 - ↳ Locks are acquired and no locks are released.
- [shrinking phase
 - ↳ Locks are released and no locks are acquired.



Lock point : Point in a transaction where the first unlock is done.

DRAWBACKS IN 2 PHASE LOCKING PROTOCOL

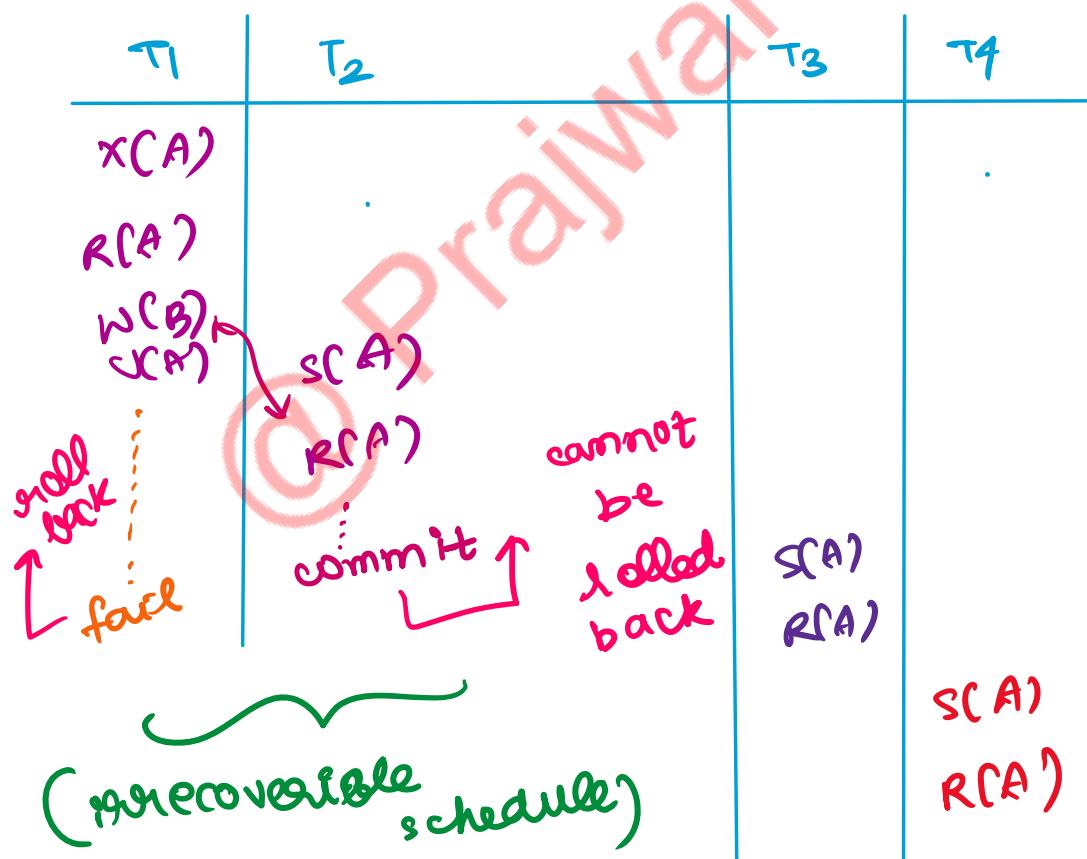
04 November 2023 18:39

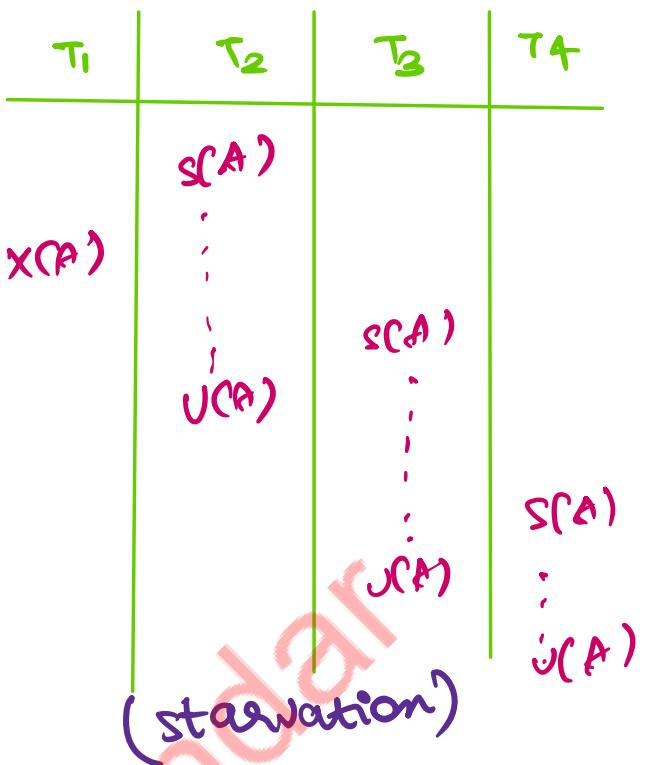
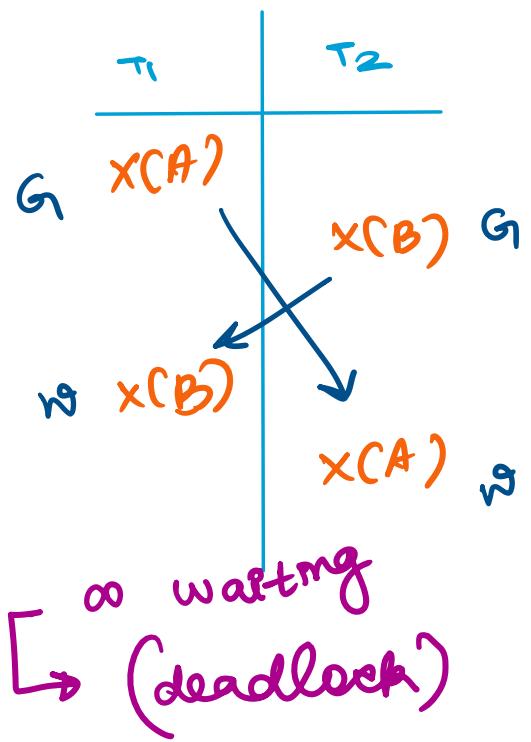
Advantages :

↳ Always ensures serializability

Disadvantages :

- ① May not be free from unrecoverability.
- ② Not free from deadlocks.
- ③ Not free from starvation.
- ④ Not free from cascading rollback.





STRICT 2PL, RIGOROUS 2PL AND CONSERVATIVE 2PL

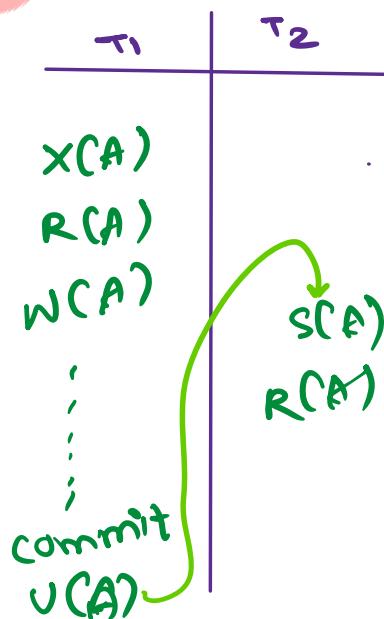
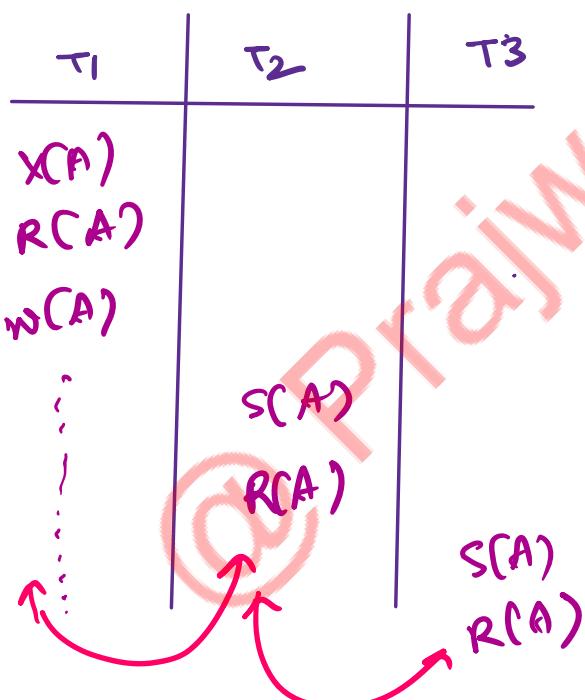
04 November 2023 19:29

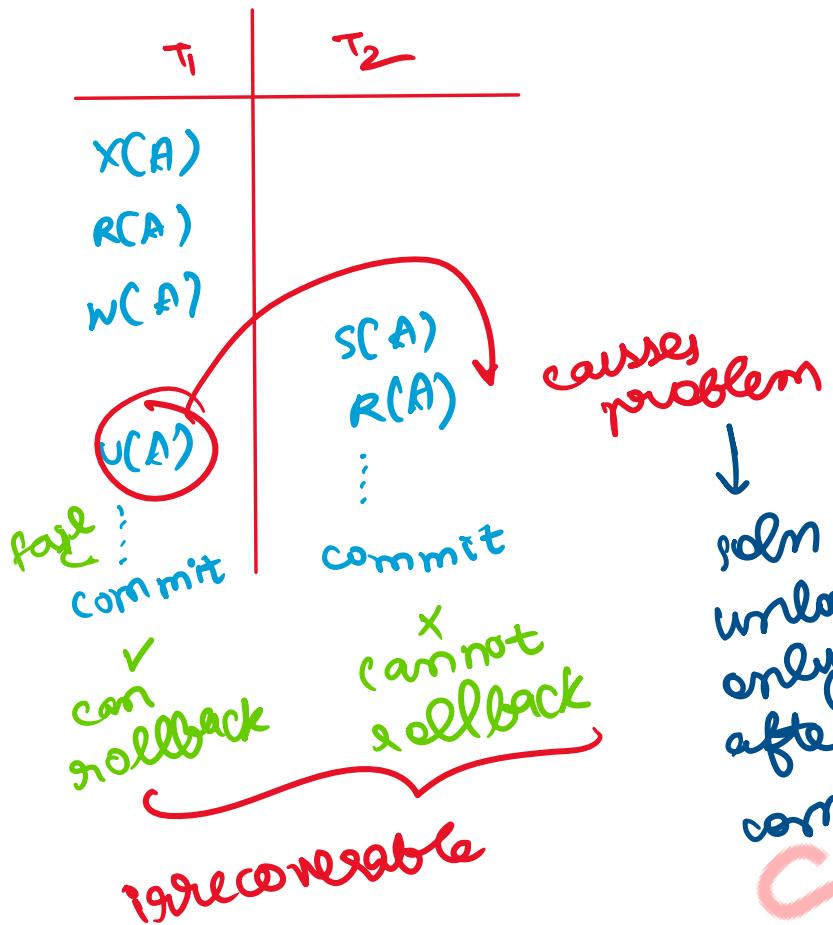
strict 2PL

- It should satisfy the basic 2PL and all exclusive locks should hold until commit / abort.

Rigorous 2PL

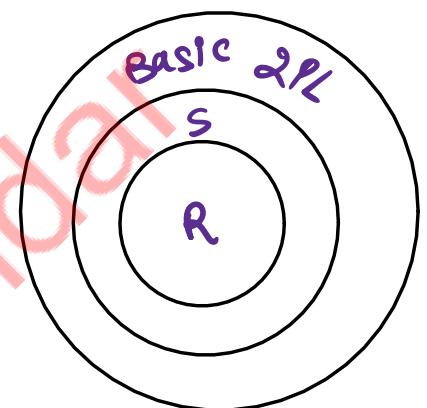
- It should satisfy the basic 2PL and all shared, exclusive locks should hold until commit / abort.



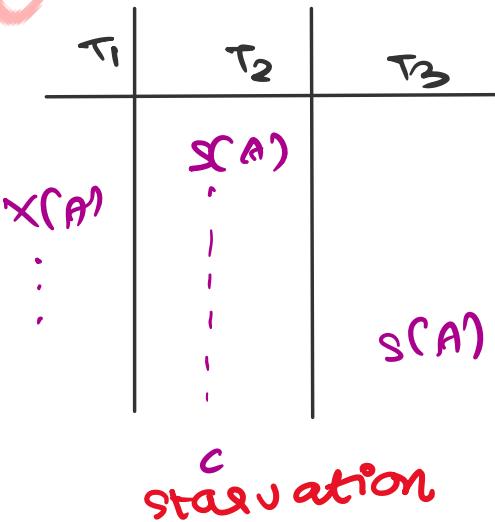
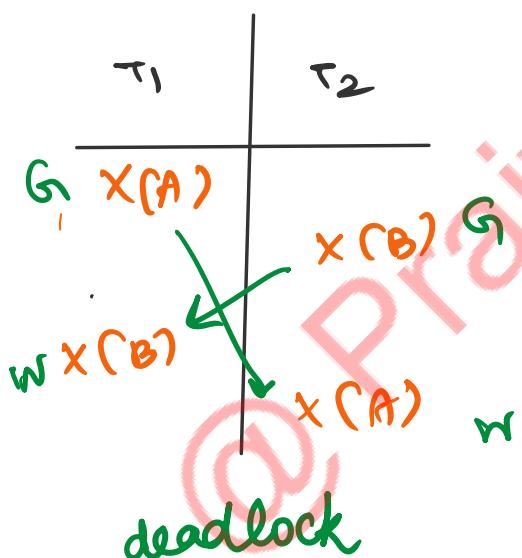


Advantages:

- ① cascadeless
- ② strict recoverable



S → strict
R → rigid



conservative 2PL

when a transaction begins, it should take all locks before starting.
[not practically possible]

TIMESNAP ORDERING PROTOCOL

04 November 2023 20:05

- unique value assigned to every transaction
- tells the order (when they enter into the system)
- $\text{read-TS (RTS)} = \text{last (latest) transaction no. which performed read successfully.}$
- $\text{write-TS (WTS)} = \text{last (latest) transaction no. which performed write successfully.}$



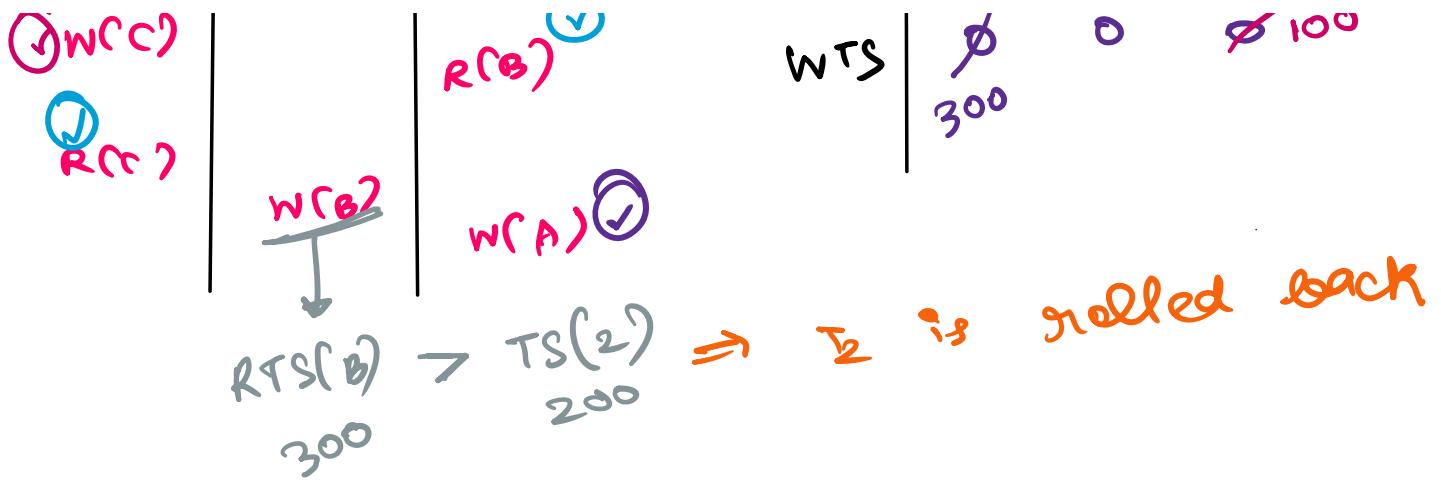
	$R(A)$	$w(A)$
	$RTS(A) = 30$	$WTS(A) = 20$

Rules :

- ① Transaction T_i issues a **Read(A)** operation
- if $WTS(A) > TS(T_i)$, rollback T_i
 - otherwise execute $R(A)$ operation
- set $RTS(A) = \max \{ RTS(A), TS(T_i) \}$
- ② Transaction T_i issues **write(A)** operation
- if $RTS(A) > TS(T_i)$, rollback T_i
 - if $WTS(A) > TS(T_i)$, rollback T_i
 - otherwise execute $w(A)$ operation
- set $WTS(A) = TS(T_i)$

T_1	T_2	T_3
✓ $R(A)$	✓ $R(B)$	
✓ $w(C)$		✓ $R(B)$

	A	B	C
RTS	100 ∅	200 ∅	300 ∅
WTS	∅ 100	∅	∅ 100



@ Prajwal Sundar