

## ARTIFICIAL INTELLIGENCE & MACHINE LEARNING

what is Artificial Intelligence?

- \* Human-like
- \* Rational-like
- \* Thinking rationally
- \* Acting rationally
- \* Acting humanly
- \* Thinking humanly
- \* Thought processes
- \* Behaviour

The Turing test

natural language knowledge representation,  
automated reasoning, machine learning,  
(vision, robotics) for full test

thinking humanly

introspection, the general problem solver,  
cognitive sciences

thinking rationally

logic, problems: how to represent and  
reason in a domain

Acting rationally

agents: perceive & act

AI examples

① common sense reasoning

- \* tweety is a bird  $\rightarrow$  it can fly
- \* Yale shooting problem
- \* smoke : fire  $\Rightarrow$  smoke

\* tweety is a bird.

All birds can fly.

Since tweety is also a bird  $\rightarrow$  it can also fly.

$\downarrow$  exception

\* tweety is an ostrich.

Ostriches cannot fly.

Since tweety is an ostrich, it cannot fly.

Yale shooting problem

Sequence of events -

Inital  $\leftarrow$  Gun is loaded (0)

$\searrow$  Alive (4)

Loaded (0)  $\rightarrow$  Loaded (1)

Alive (4)  $\rightarrow$  Alive (3)

$\downarrow$

like a search problem } used typically in games, move from one state to another state till the goal is reached.

	loaded	alive
{ many problem states }	0	4
	1	4
	1	3
	2	2

Final goal: To kill all

loaded  $\rightarrow$  anything

alive  $\rightarrow$  0

② Update or revise knowledge :

$$A \text{ cost } B \rightarrow C$$

observe  $c = 0$ , vs Do  $c = 0$

③ chaining theories of actions

- { Looks like P → P is P
- Make looks like P → Looks like P
- Make looks like P → P is P ??

eg

- { Birds → fly
- { Tweety → bird
- Tweety → fly

- \* Garage door example : garage door not included.
- \* planning benchmarks .
- \* & puzzle, & Queen, Block world,  
Grid - maze world
- \* Abduction : Cambridge parking example.

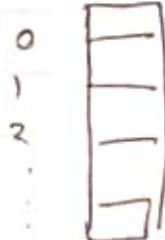
PROGRAM : I  $\Rightarrow$  TIC TAC TOE

data structure approach

1	2	3
4	5	6
7	8	9

$\Rightarrow 3^9$  combinations  
[empty, X, O]

$$3^9 \rightarrow$$



unnecessary space , looking  
at all combinations  $\rightarrow$  a  
waste of time

Another DS approach

2 → blank    3 → X    5 → 0

9x9 vector     $\langle 1, 2, \dots, 9 \rangle$

$\langle 2, 3, 5, \dots \rangle$

\*

returns 5 if centre is blank

using prefix and suffix

goes function → either help produce  
a win (or) draw

state space tree

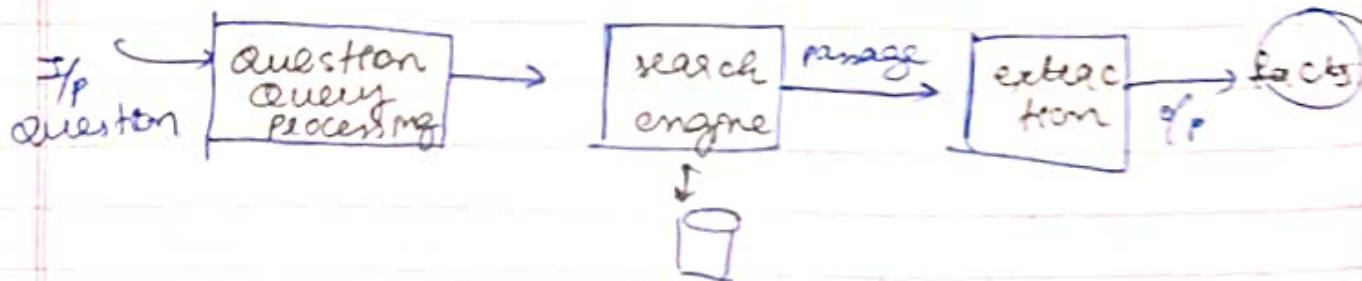
$\alpha - \beta$

throw unnecessary nodes & keep only  
those which are useful

e.g. automatic draw / resign by the  
computer during chess  $\Rightarrow \alpha - \beta$  algorithm

Q - A system

Question ↗  
 ↙ definitional search, taking a  
 ↙ message & giving  
 ↙ Factorial ↗  
 ↙ facts, intelligence needed



e.g. bengalor given

## AGENTS (CHAPTER-2)

contents :

- \* Agent and environments
- \* Rationality
- \* LPEAS ( performance measure, environment, actuators, sensors )
- \* Environment types
- \* Agent types

Agents :

An agent is anything that can be viewed as perceiving its environment through sensors and acting upon that environment through actuators.

Human agent

sensors : eyes, ears, other organs

Actuators : hands, legs, mouth, other body parts

Robot agent

sensors : cameras & infrared range finders

actuators : various motors

• Agents and environments

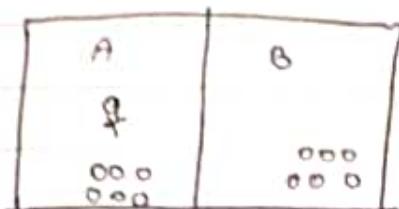
Agent function, maps from percept

the states to actions.  $f : P^+ \rightarrow A$

Agent program : runs on the physical architecture to produce f

agent = architecture + program

## VACUUM-CLEANER WORLD



Percepts : location and content, eg [A, dirty]

Action: left, right, suck, noop

## Rational Agents



An agent should strive to "do the right thing" based on what it can perceive and the actions it can perform. The right action is the one that will cause the agent to be most successful.

Performance measure: an objective criterion for success of an agent's behaviour.

eg: performance measure of a vacuum-cleaner agent could be:

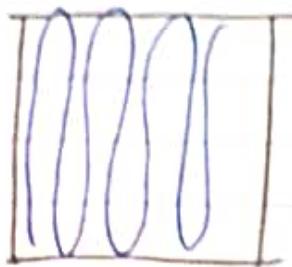
- \* amount of dirt cleaned up
- \* amount of time taken
- \* amount of electricity consumed
- \* amount of noise generated, etc

$$P \rightarrow f(D, T, E, N, M) \hookrightarrow (\text{no. of money})$$

## work of Rational Agent

For each possible percept sequence, a rational agent should select an action that is expected to maximize its performance measure, given the evidence

provided by the percept sequence and whatever built-in knowledge the agent has.



what is involved in intelligence?

Intelligent Agents,

- Ability to interact with the real world
- knowledge representation, reasoning and planning
- learning and adaptation

usually challenged ATM  
face identification + PIN

80% → data for training  
20% → data for testing (unseen data)

implementing agents

- Table look-ups
- Autonomy
- structure of an agent  
agent = architecture + program

Autonomy

↳ even in a new environment, the agent should be in a position to design an act

## omnipotence + learning + autonomy

### omnipotence

- ↳ actual outcome of an action that is carried out

vacuum cleaner → 2 cameras

algorithm → camera captured info → convert into a rectangle of coordinates

### Reflex vacuum Agent

function REFLEX-VACUUM-AGENT

([location, status]) returns an action

if status = dirty then return Suck

else if location = A then return Right

else if location = B then return Left

function TABLE-DRIVEN (receipt) returns an action

table  
driven  
agent

receipt sequence → empty

table → table of actions

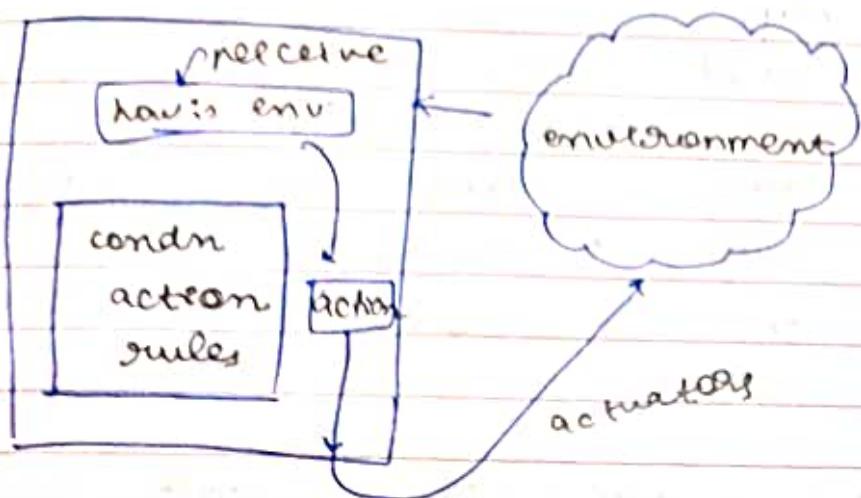
indexed by receipt sequence

append (receipt) to the end of receipt

action ← look-up (receipt, table)

### types of agents

- simple reflex agent
- model based reflex agent
- goal based agent
- utility based agent

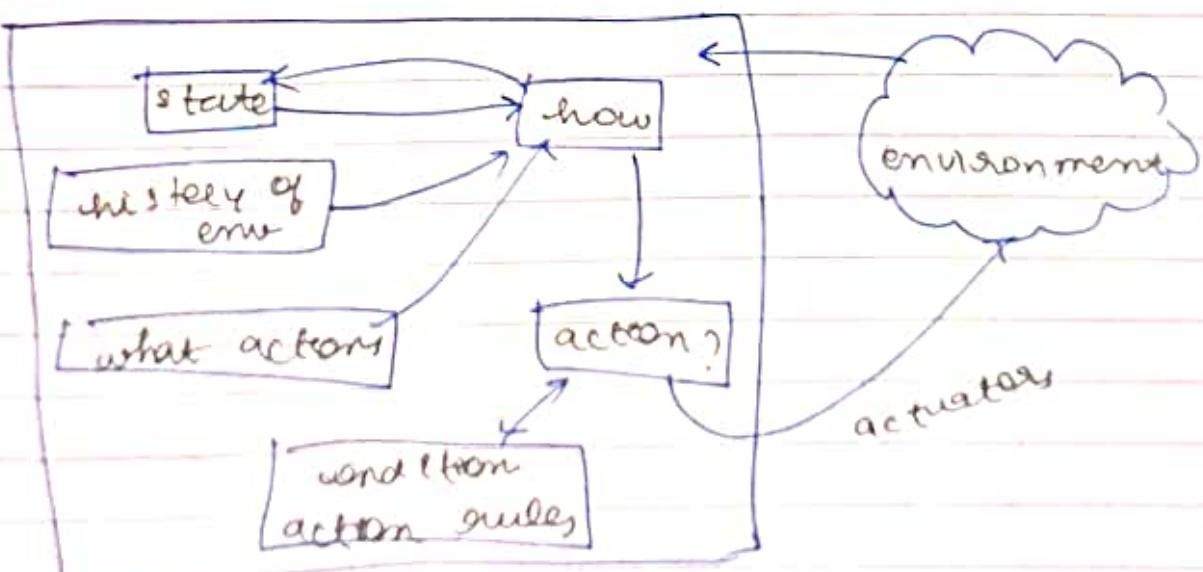


Automated driving car

red light  $\rightarrow$  brake  
 $(\text{near vehicle}, \text{red}) \rightarrow \text{break}$

$(\text{near vehicle}, \text{no red}) \rightarrow \text{drive}$

Model-Based Reflex agent



maintaining & updating history as an when we try to capture

## Google maps

model - based reflex agent (percept),  
returns action

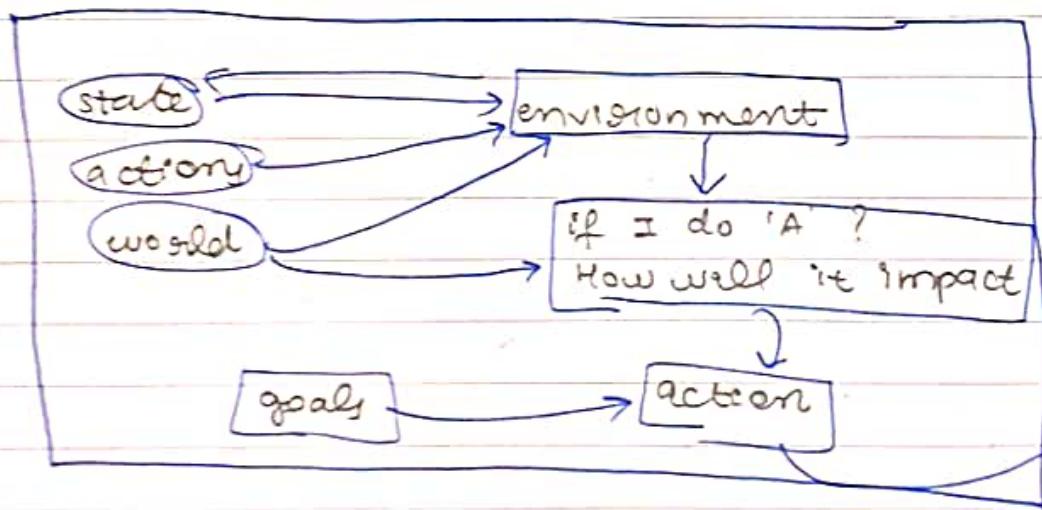
- state → current environment
- model → description of the state
- rules → conditions & actions
- actions → to be carried out

state ← update - state (state, action,  
percept, model)

Rule ← rule - match (state, rules)

Action ← rule action

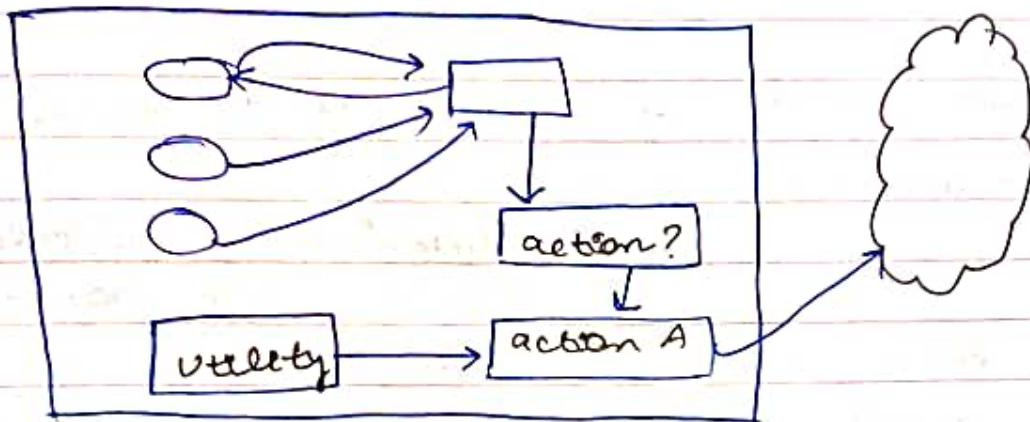
## Goal-based Agent



Goal-based agents are modifications to model-based agents which add goals to the parameters and deciding what action to be carried out.

Automated driver : red, green, change signals, no overtake, keep left (or right), etc.

utility - Based Agent → user experience  
 \* comfortable  
 \* happy  
 goal / model /  
 experience



goal - based agents & utility - based agents work by a lot of searching & planning.

### TYPES OF ENVIRONMENT

① Fully observable vs partially observable

↓  
complete environment  
is available to us

↓  
only a portion of the  
environment is  
available to us.

Unobservable environment

we are not able to get the environment  
at all

Google - maps : fully observable,  
some information → return action  
senses the complete environment at  
any given point of time

Autonomous vehicles : also fully observable  
if cameras are kept all over

Vacuum cleaner : partially observable environment

② single-agent vs multi-agent (Drivers)  
(crossword puzzle)

autonomous vehicle  
driver → one agent

other drivers → also agents with which  
driver should interact to make a decision

single-agent → one agent

more than one agent → multi-agent

③ Deterministic vs stochastic.

↓  
next state depends  
upon the current  
state & action

↑  
next state does not  
depend upon the current  
state & action

Stochastic → thought of as a partially  
observable environment  
similarly Deterministic ↔ fully observable

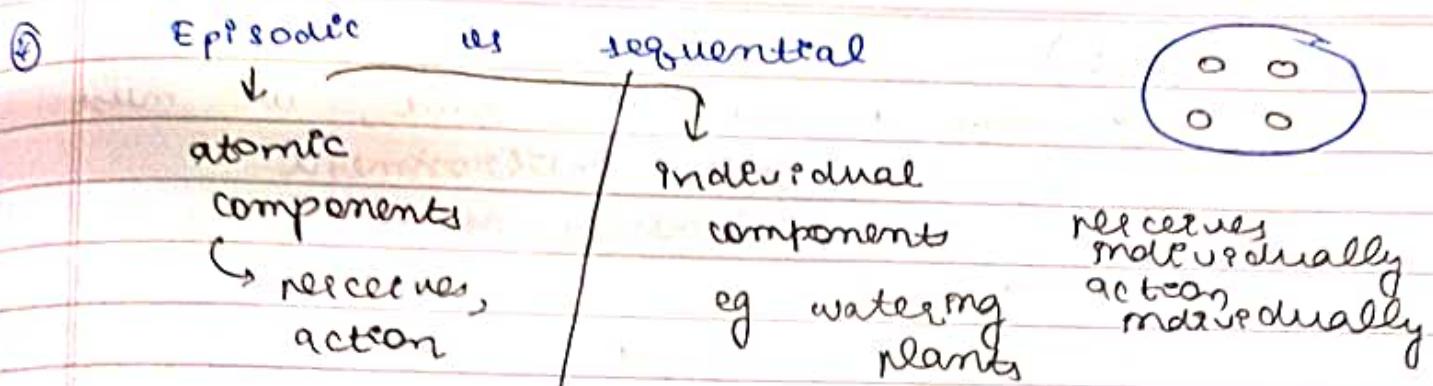
Taxi - driving : multi-agent environment

/  
deterministic

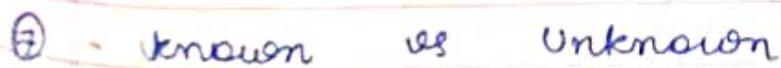
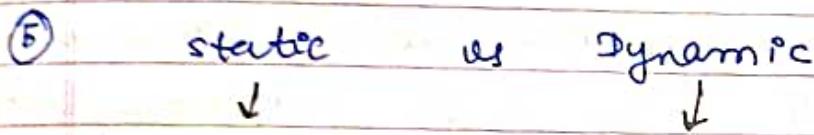
\  
stochastic

perfect driver

✗ (not always)



current action depends upon  
the previous action



## PEAS

To design a rational agent, we must specify the task environment.

P → Performance measure

E → Environment

A → Actuators

S → Sensors

e.g. taxi (automated)

P → safety, destination, profit, legality, comfort

E → US streets/freeways, traffic, pedestrians, weather

A → steering, accelerator, brake, horn, speaker, dest

S → video, accelerometers, gauges, engine sensors, keyboard, GPS.

## Environment types

	NP	solitaire	Tutor	Backgammon	Internet shopping	Taxi
partially observable	✓	✗	partially	✓	✗	✗ partially
deterministic	✗	✓	✗	✗	partially	✗
Episodic	✗	✗	✗	✗	✗	+
static	✗	✓	✗	semi	semi	+
stochastic	✗	✓	✓	✓	✓	+
single-agent	✓	✓	✗	✗	no except auctions	+

the environment type largely determine the agent design.

The real world is : partially observable, stochastic, sequential, dynamic, continuous and multi-agent

MD → medical diagnosis

## PEAS : English tutor

- P : student's performance  $\Rightarrow$  oral written
- E : set of students, testing agency / test
- A : exercises, suggestions, speech, corrections
- S : keyboard only / oral english capture

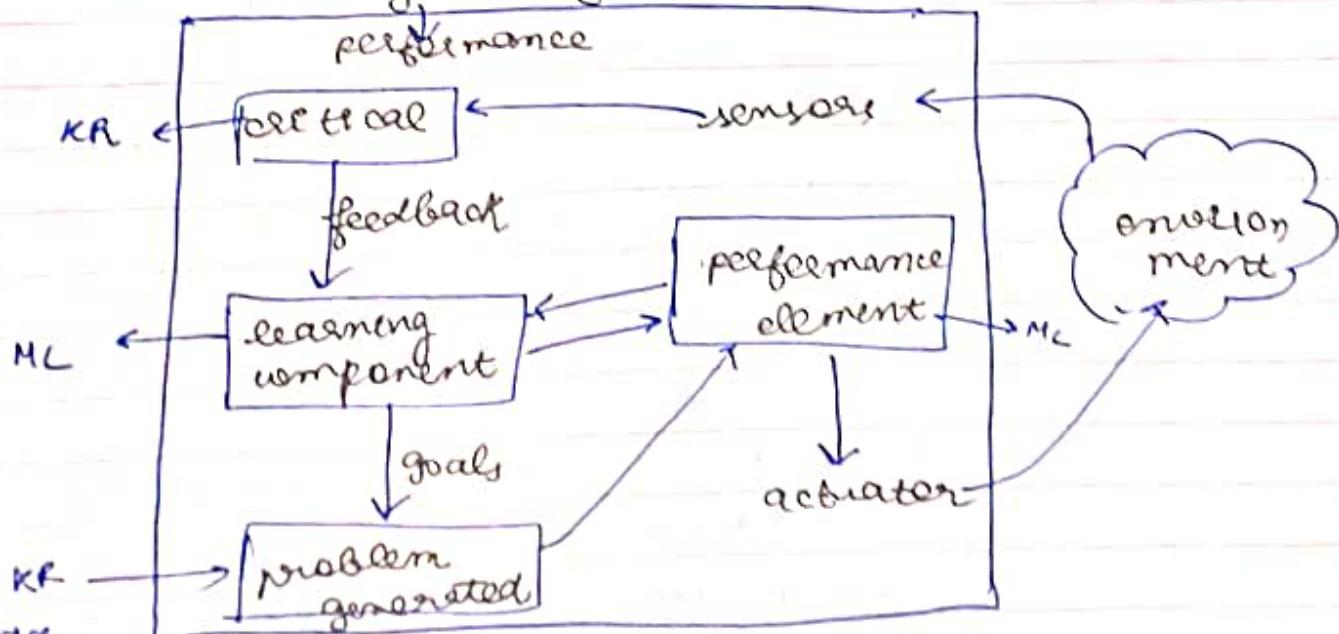
## PEAS : medical diagnosis

↓

- P : healthy patient, low cost
- E : hospital, patient, staff, lab
- A : diagnostic report, treatment, referral, O/T
- S : history of symptoms, patient's answers

learning agent

↳ intelligent agent



knowledge engineering

Learning element → make modifications/  
outcome better

Performance element → actual action

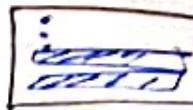
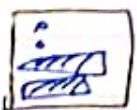
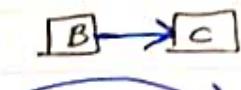
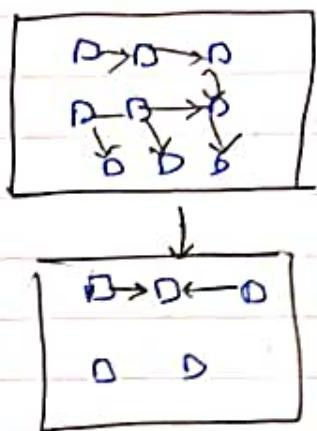
ML → machine learning

KR → knowledge representation

KRn → knowledge reasoning

### Environment components :

- ① Atomic
- ② factored
- ③ structured



} split each state into  
} a set of fixed  
variables / attributes

} solved separately  
combine to get overall sol.

objects  
attributes

relationships among objects

### Examples :

atomic → game / Hidden Markov models

HMM  
Markov Decision Process  
MDP

factored → CSP constraint satisfaction problem  
propositional logic /  
Bayesian networks, ML

structured → First order logic / FOL  
natural language understanding / NLP  
knowledge based understanding

knowledge representation KR → ontology →  
entity & their relationships

## Planning Agents :-

- "what if"
- decisions based on (hypothetical) consequences of actions
- must have a model of how the world evolves in response to actions
- must formulate a goal (test)
- consider how the world would be

Optimal vs complete planning

Planning vs replanning

search problems

problem solving agents

### • Atomic representation

states are considered as whole with no internal structure visible to the problem

### • Goal formulation

Based on the current situation and the agent's performance measure

### • Problem formulation

what actions and states to consider given a goal.

### • Environment - observable, discrete, deterministic.

search - process of looking for a sequence of actions  
takes problem as input & returns solution  
in the form of an action  
actions → carried out in the execution phase

- \* Initial state - starting point
- \* Actions - description of the possible actions available
- \* Applicable actions - From a state ' $s$ ' what actions are possible
- \* Transition model - A description of what each action does
- \* Successor - states reachable from a given state
- \* State space - set of all states reachable from the initial state by any sequence of actions
- \* Graph - Nodes are states & links between the nodes are actions
- \* Path - sequence of states

*transition model*  $\left\{ \begin{array}{l} s \rightarrow \text{action } 'a' \rightarrow p \\ s(a) = p \end{array} \right.$

- \* Goal test - test to verify whether a state is a goal
- \* Path cost - function to assign a numeric cost to each path
- \* Solution - leads from start state to final state

### Search Problem

A search problem consists of

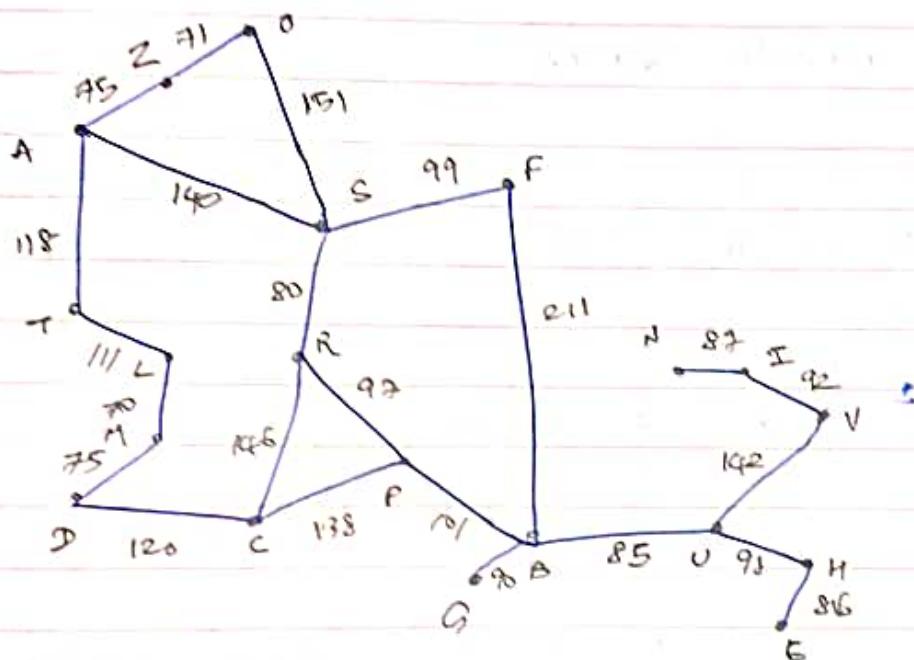
- \* state space
- \* successor function  
(with actions, costs)
- \* start state & goal test

Search problems are modeled

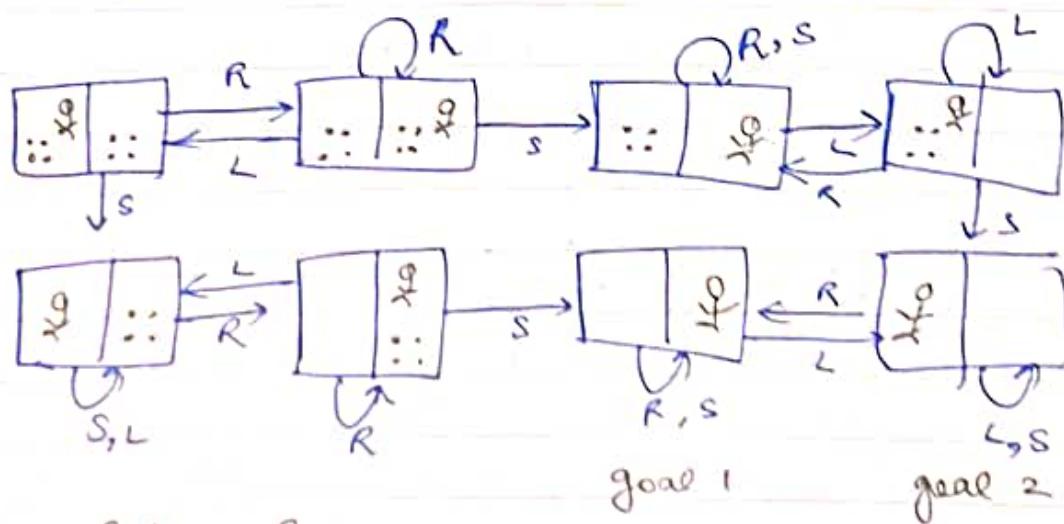
## Example : Travelling in Romania

state space  $\rightarrow$  citiessuccessor function  $\rightarrow$  roads : go to an adjacent city with cost = distance

start state : A

Goal test : Is state = B ?  
solution ?

state space for a vacuum world


 $\left\{ \begin{array}{l} L \rightarrow \text{left} \\ R \rightarrow \text{right} \\ S \rightarrow \text{suck} \end{array} \right\}$

$n$  locations :  $n \times e^n$  states

Here  $n=2$ ,  $2 \times e^2 \Rightarrow$  agent  $\&$  location states

Initial state :

Actions : L, R, S

Transition Model :

Goal test : all squares are clean

Path cost : each step cost 2  $\rightarrow$  no steps

## 8 puzzle Game

7 2 4		1 2 3	left, right,
5	6	4 5 6	up, down
8 3 1		7 8	movement

9 locations  $\Rightarrow (9 \times e^9$  states)

## 8 Queens problem

place 8 queens on the chessboard such that no 2 queens attack each other

## Route finding

states : locations

Initial state : user ?

Actions : mode of transport (run, walk, car, auto, etc)

Transition model : city / place  $\rightarrow$  place

Goal test : reached final destination

Path cost : calories / petrol

state space :

world state

↳ includes every last detail of the environment

search state

↳ keeps only the details needed for planning (abstraction)

problem : pathing

states :  $(x, y)$  location

Actions : N S E W

successor : update location only

Goal test :  $\exists (x, y) = \text{END}$

Problem : Eat-All-Dots

states :  $\{(x, y), \text{dot booleans}\}$

Actions : N S E W

successor : update location & possibility  
a dot boolean

state space Graph

- ★ mathematical representation of a search problem
- ★ nodes are abstracted world configurations
- ★ arcs represent successors (action results)
- ★ goal test is a set of goal nodes  
(may be only one)

- ★ In a state space, each state occurs only once
- ★ Full graph in memory  $\rightarrow$  too big, but useful

## search tree

A "what if" tree of plans & outcomes

start state  $\rightarrow$  root node

children  $\rightarrow$  successors

nodes show states  $\rightarrow$  correspond to PIA

& achieve those states

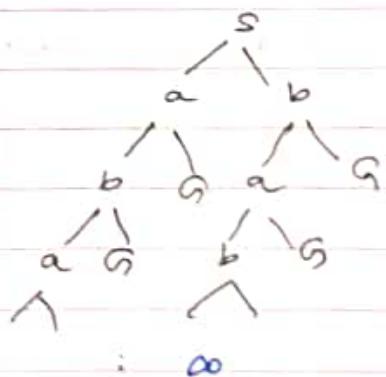
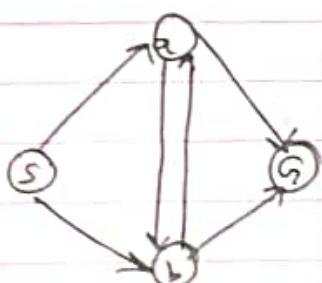
we never actually build the whole tree

we construct nodes & paths on demand

& we construct as little as possible

$$\begin{array}{ccc} \text{NODE} & \longleftrightarrow & \text{PATH} \\ (\text{search tree}) & & (\text{state space graph}) \end{array}$$

state space graphs as search trees,



(lots of repeated structures in the search tree)

$\downarrow$   
construct only  
these needed

- searching with a search tree

search

- \* expand out potential plans (tree nodes)
- \* maintain a fringe of partial plans under consideration
- \* try to expand as few tree nodes as possible

## General tree search

function TREE-SEARCH (problem, strategy)  
    returns a solution, or failure

    initialize the search tree using the initial state of problem

    loop do

        if there are no candidates for expansion  
            then return failure

        choose a leaf node for expansion  
            according to strategy

        if node contains a goal state then  
            return the corresponding solution  
        else expand the node and add the  
            resulting nodes to the search tree

    end

important ideas :

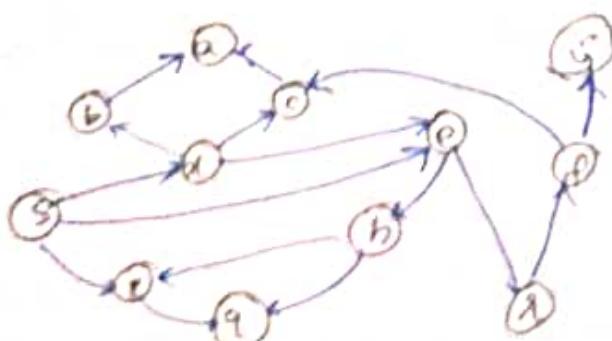
- fringe
- expansion
- exploration strategy

main Q: which fringe node to explore?

## depth-first search

strategy : expand a deepest node first

implementation : Fringe is a LIFO stack



## search algorithm properties :

complete : Guaranteed to find a solution if one exists ?

optimal : Guaranteed to find the least cost path

time complexity ?

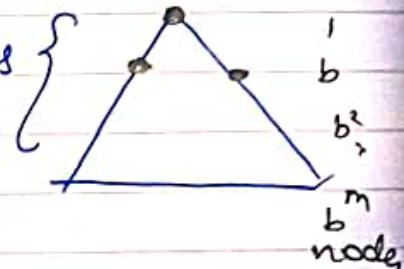
space complexity ?

cartoon of search tree :

$b \rightarrow$  branching factor

$m \rightarrow$  maximum depth

solutions at various depths



no. of nodes in entire tree

$$1 + b + b^2 + \dots + b^m = O(b^m)$$

## depth-first search (DFS) properties

strategy : expand a deepest node first

what nodes DFS expands ?

\* some left prefix of the tree.

\* could process the whole tree

\* If  $m$  is finite, takes time  $O(b^m)$ .

How much space does the fringe take ?

only has siblings on path to root so  $O(M)$ .

Is it complete ?

$m$  could be infinite, so only if we prevent cycles (more later)

Is it optimal ?

No, it finds the "leftmost" solution regardless of depth or cost.

## Uninformed search techniques

↓  
ahead of time  
goal state is not known, if the goal is hit while we keep expanding, it is done.

- ① DFS (Depth First search)
- ② BFS (Breadth First search)
- ③ Uniform cost search
- ④ Depth Limited search
- ⑤ Iterative deep search
- ⑥ Bidirectional search

completeness ✓ optimality ✓  
time complexity ✓ space complexity ✓

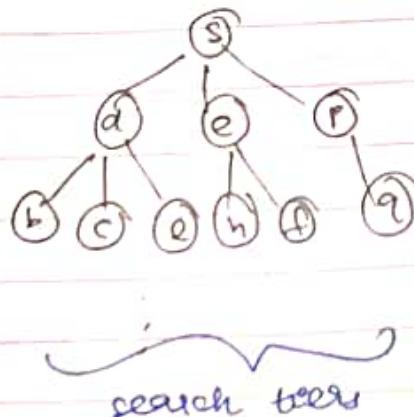
### Depth Limited search

↳ no guaranteed solution [ same as in DFS]  
but it limits the branching factor  
to d depth, prefixed by the user.

### Breadth First search

strategy: expand a shallowest node first

Implementation: fringe is a FIFO Queue.



queue  
S d e p b c e h r g

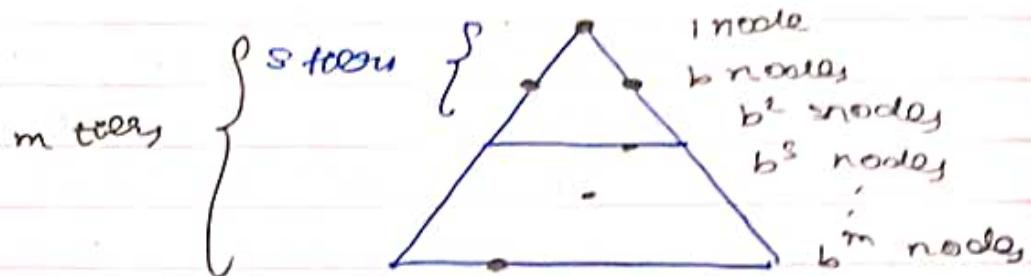
[ same graph as used  
in DFS ]

S → e → r → f → G

## Breadth First search Properties

what nodes does BFS expand?

- ★ processes all nodes above shallowest solution
- ★ let depth of shallowest solution be  $s$
- ★ search time takes  $O(b^s)$



How much space does the fringe take?

Has roughly the last tree, so  $O(b^s)$ 

Is it complete?

$S$  must be finite if a solution exists  
so yes!

Is it optimal?

only if costs are all 1 (more on costs later)

when will BFS outperform DFS?

Goal state with less connected graph  $\rightarrow$  BFS outperforms DFS.

when will DFS outperform BFS?

complete graph & every node is connected to every other node in all possible scenarios.

leftmost solution is obtained faster in DFS,  
outperforms BFS

## Iterative Deepening

Idea: get DFS's space advantage with BFS's time / shallow - solution advantage

• Run a DFS with depth limit 1. If no soln:  
increases limit: 0, 1, 2 limit  $\rightarrow$  2

- \* Run a DFS with depth limit 2. If no soln:
- \* Run a DFS with depth limit 3.

Isn't that wastefully redundant?

Generally most work happens on the lowest level searched, so not so bad!

## cost-sensitive search

BFS finds the shortest path in terms of number of actions. It does not find the least cost path. Another similar algorithm helps find the same.

## uniform cost search

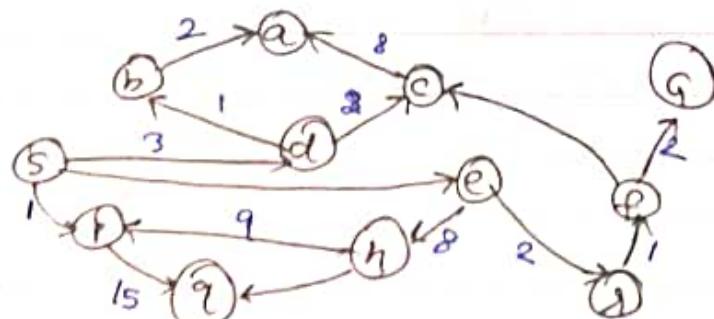
→ variation of BFS which gets the least possible cost.

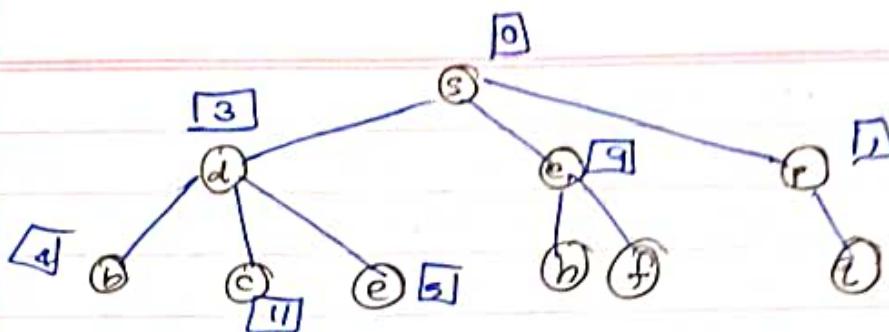
Fringe  $\Rightarrow$  greedy in nature.

put in queue based on cost

strategy: expand a cheapest node first

Fringe is a priority queue (priority: cumulative cost)





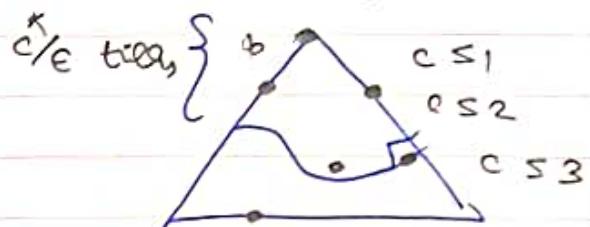
In priority queue: ~~stackable~~  
3 9 10

keep removing, adding and updating  
values from the queue.

uniform cost search (ucs) properties

what nodes does ucs expand?

- \* Processes all nodes with cost less than cheapest solution!
- \* If that solution cost  $c^*$  and avg cost at last  $e$ , then the "effective depth" is roughly  $c^*/e$ .



How much space does the fringe take?

Has roughly the last tier, so  $O(b^{c^*/e})$   
Is it complete?

Assuming best solution has a finite  
cost and minimum arc cost is  
possible, yes!

Is it optimal?

Yes! (proof via A\*)

## uniform cost issue,

UCS explores increasing cost contours  
the good: UCS is complete & optimal  
the bad:

- \* explores options in every "direction"
- \* no information about goal location

## The one Queue

All these search algorithms are the same except for fringe strategies.

& conceptually, all fringes are priority queues (i.e. collections of nodes with attached priorities).

- \* Practically for DFS & BFS you can avoid the leg(n) overhead from an actual priority queue by using stacks & queues.
- \* can even code one implementation that takes a variable queuing object

## search and Models

search operates over models of the world

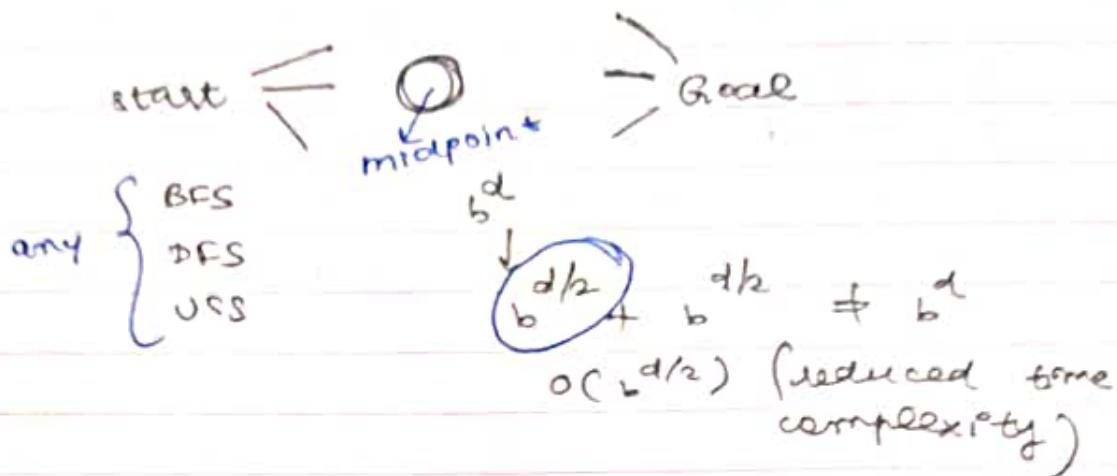
- & the agent doesn't actually try all the plans out in the real world!
- \* planning is all "in simulation"
- & your search is only good as your model

## Example: Pancake problem

stack of pancakes, flip them such that smallest pancake is on top & all of them are in increasing order

(flip)  
  
 [scoring by prefix reversal] How many flips are needed to finally achieve this cost: no. of pancakes flipped

## Bidirectional search



## Informed search



Goal is already known  $\rightarrow$  more information is available about it than in the uninformed searches.



(Heuristic function)

## Search Heuristic

General approach: best first search

which node to expand first?

Tree search / Graph search

$f(n) \rightarrow$  evaluation function

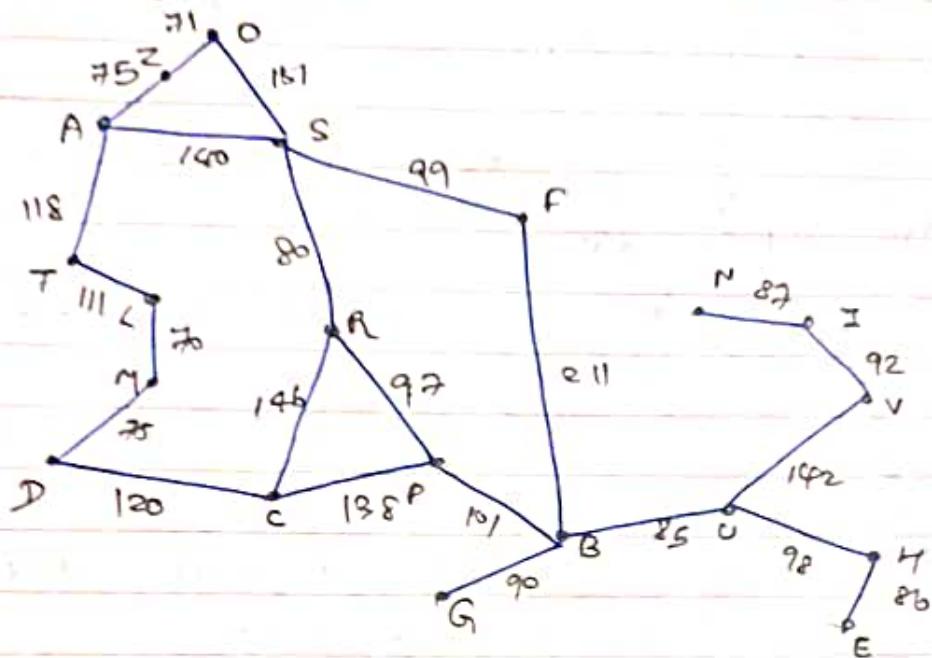
$f(n) \Rightarrow$  heuristic function

$f(n) =$  cost of lowest path from start node ' $s$ ' to goal state ' $g$ '

A heuristic is :

- \* A function that estimates how close a state is to a goal
- \* designed for a particular search problem e.g. manhattan distance, euclidean distance for pathfinding

$\hookrightarrow$  SLD : straight line distance  
 $\hookrightarrow$  eg for a heuristic function



straight line distances to B :

$g(x)$

A - 366	B - 0	C - 160	D - 242
E - 161	F - 178	G - 77	H - 151
I - 226	J - 244	K - 241	L - 234
M - 241	N - 234	Q - 193	R - 253
O - 380	P - 98	V - 199	W - 344
S - 329	U - 80	Z - 344	
T - 329			

lat, lon  $\rightarrow$  Haversine distances  $\rightarrow$  used to estimate SLD between 2 points

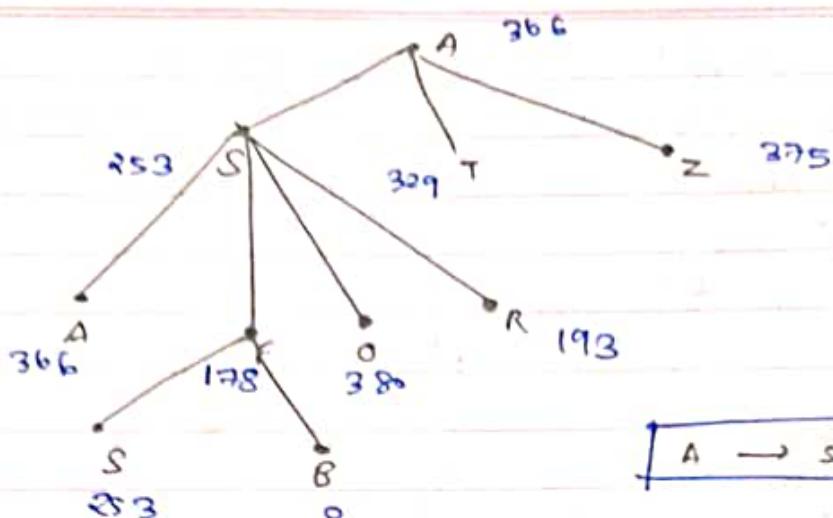
For pancake problem

Heuristic : the no. of the largest pancake that is still out of place

Greedy search

Greedy best first search  
 $\hookrightarrow$  SLD  $\Rightarrow h(x)$

SLD is used as a basis for expansion



(greedy approach  
↓  
SLP)

$$A \rightarrow S \rightarrow F \rightarrow B$$

$O(n^m)$  worst case time & space complexity

### Greedy search

strategy: expand a node you think is closest to a goal state

Heuristic: estimate of distance to nearest goal for each state

A common case:

Best-first takes you straight to the (wrong) goal

worst case: like a badly guided DFS

### A\* search

variations: memory bounded  $A^* \Rightarrow M A^*$   
(Greedy best)  
simplified memory bounded  $A^* \Rightarrow S M A^*$

In GBFS it finds path from new state to goal state, but path from the source to new state was not controlled.

$A^*$  search considers this aspect also

- $g(n) \rightarrow$  cost to reach a node
- $h(n) \rightarrow$  cost to get from the node to goal

$$f(n) = g(n) + h(n)$$

$$f(n) = g(n) + h(n)$$

$g(n) \rightarrow$  path cost from start to node  
 $h(n) \rightarrow$  sld (if  $f(n) \rightarrow$  heuristic)

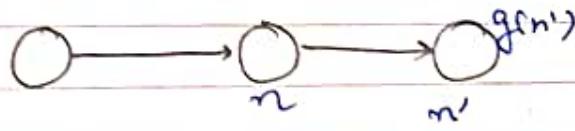
- $A^*$   $\Rightarrow$  UCS  $\Rightarrow (g+h)$  instead of  $g$
- Admissibility
- consistency

### Admissibility

admissible heuristic

not over-estimate the cost to reach the goal

$$g(n) = n \quad f(n) = g(n) + h(n)$$



consistency  $\downarrow$  stronger condition

monotonicity

$A^* \rightarrow$  graph search

$h(n)$  is consistent if

$\forall n$ , and  $n'$  is a successor of  $n$ , the estimated cost of  $n'$  is not greater than the

step cost of getting to  $n'$  from the goal  $\leftarrow n'$

### Triangular inequality

$$h(n) \leq c(n, a, n') + h(n')$$

$c(n) \rightarrow$  cost to reach the goal from  $n$

$h(n') \rightarrow$  cost to reach the goal from  $n'$

correctness



300km



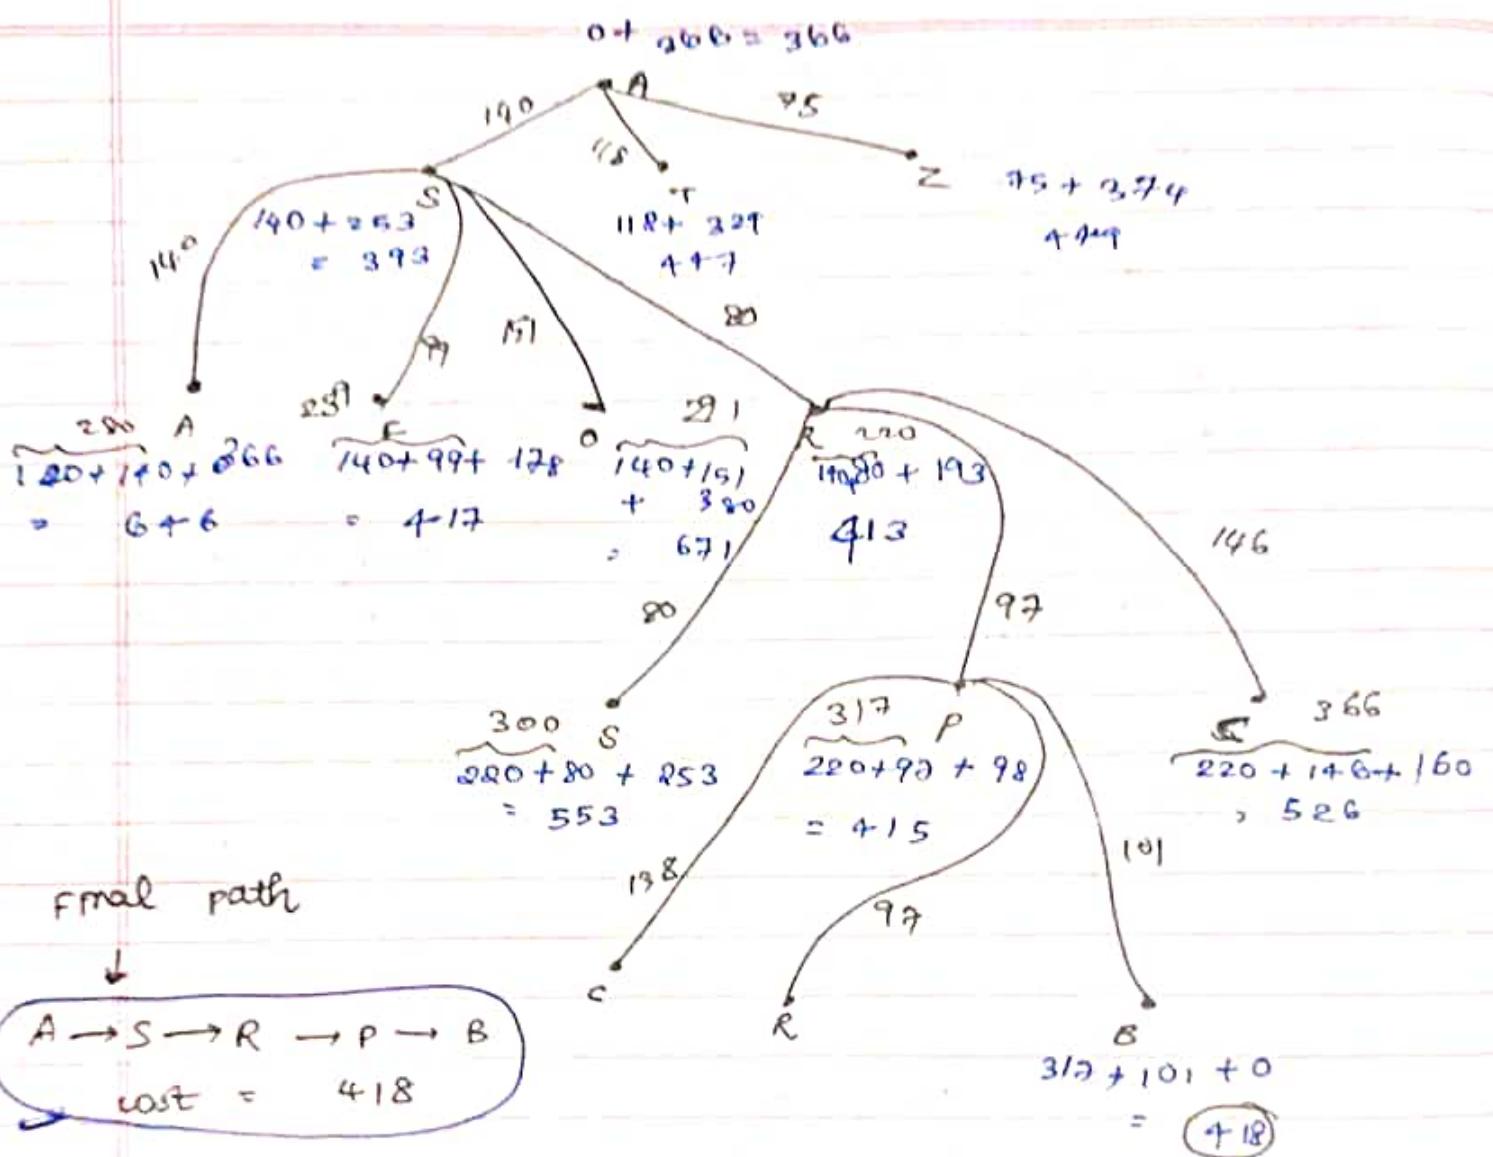
middle

250km

$\text{O } x \text{ (must be } > 300\text{km})$

Every consistent heuristic is admissible.

$\downarrow$   
 follows the triangular inequality

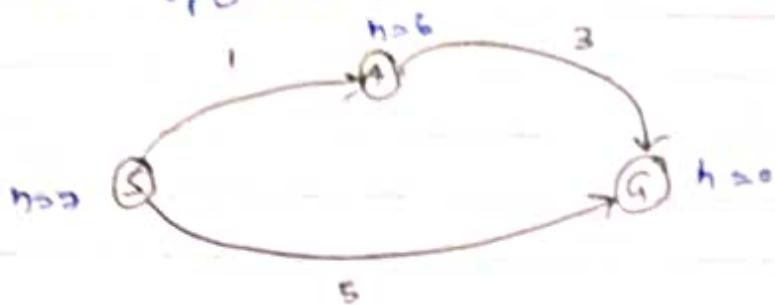


- \* uniform cost orders by path cost or backward cost  $g(n)$
- \* Greedy Orders by goal proximity or forward cost  $h(n)$
- \* A\* search orders by the sum  $f(n) = g(n) + h(n)$

when should A\* terminate?

should we stop when we enqueue a goal? No, only stop when we dequeue a goal.

Is A\* optimal?



what went wrong?

Actual bad goal  
cost < estimated  
good goal cost

Idea: Admissibility

Inadmissible (pessimistic) heuristics break optimality by trapping good plans on the fringe.

Admissible (optimistic) heuristics slow down bad plans but never outweigh true costs

Admissible Heuristics

A heuristic  $h$  is admissible (optimistic) if  $0 \leq h(n) \leq h^*(n)$

where  $h^*(n)$  is the true cost to a nearest goal

$$h(n) \leq c(n, g, n^*) + a(n^*)$$

optimality of A\* Tree search

Assume:

- A is an optimal goal node
- B is a suboptimal goal node
- $h$  is admissible



claim: A will exit the fringe before B

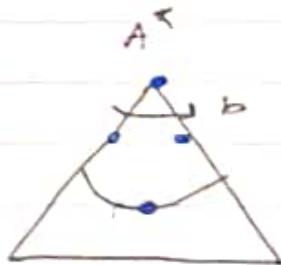
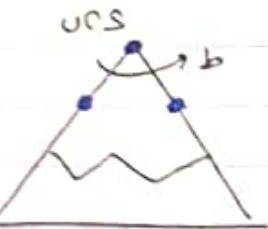
optimality of A\* Tree search: Blocking

Imagine B is on the fringe  
some ancestor of A is on the fringe too  
(maybe A)

- claim:  $n$  will be expanded before  $B$
- ①  $f(n) \leq f(A)$   
 $f(n) = g(n) + h(n)$  defn of f-cost  
 $f(n) \leq g(A)$  admissibility of  $h$
  - ②  $f(A) < f(B)$   
 $g(A) < g(B)$   $B$  is suboptimal  
 $f(A) < f(B)$   $h=0$  at goal
  - ③  $n$  expands before  $B$

All ancestors of  $A$  expand before  $B$   
 $A$  expands before  $B$

### Properties of $A^*$



UCS vs  $A^*$  contours

UCS expands equally in all directions



$A^*$  expands mainly toward the goal,  
but does hedge its bets to ensure optimality



### A\* Applications

- \* video games
- \* pathing / routing problems
- \* resource planning problems
- \* robot motion planning
- \* language analysis
- \* machine translation
- \* speech recognition

Creating Admissible Heuristics

↳ (most work)

admissible heuristics are solutions to relaxed problems, where new actions are available

[ sometimes inadmissible heuristics are useful too ]

Example : 8 puzzle

what are the states?

how many states?

what are the actions?

how many successors from the start state?

what should the cost be?

heuristic : no. of tiles misplaced.

$h(\text{start}) = 8 \Rightarrow$  relaxed problem heuristic

[ what if we had an easier 8 puzzle where any tile could slide any direction at any time pushing other tiles? ]

Total Manhattan distance / city block distance  
sum of all edge costs

using actual cost as a heuristic?

Admissible?

sane on nodes expanded?

what wrong with it?

$A^*$   $\Rightarrow$  trade off between quality of estimate and work per node

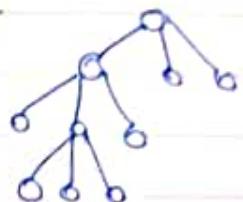
Memory Bounded Heuristic search

Iterative deepening -  $A^*$  ( $IDA^*$ )

$$IDA^* \rightarrow ID$$

$$h(n) = f(n) + g(n)$$

↓



reduces amount of memory needed to store the intermediate results

Recursive Best First search

RBFS  $\rightarrow$  linear search  $\rightarrow A^*$  modification  
similar to recursive depth first search  
limit + alternate paths are also remembered

RBFS algorithm

(problem) returns a soln / failure  
return RBFS (problem, make-node (problem,  
initial state),  $\infty$ )

RBFS (problem, node, f-limit)

returns soln / failure

If (problem goal-test (node.state)) return solution  
successor [ ],

for each action in problem.action (node.state)  
add-child-node (problem, node, action) to  
successor

if successor is empty return failure

& 's' in successor

$$s.f = \max(s.g + s.h, \text{node}.f)$$

do

best = lowest f-value node in successor

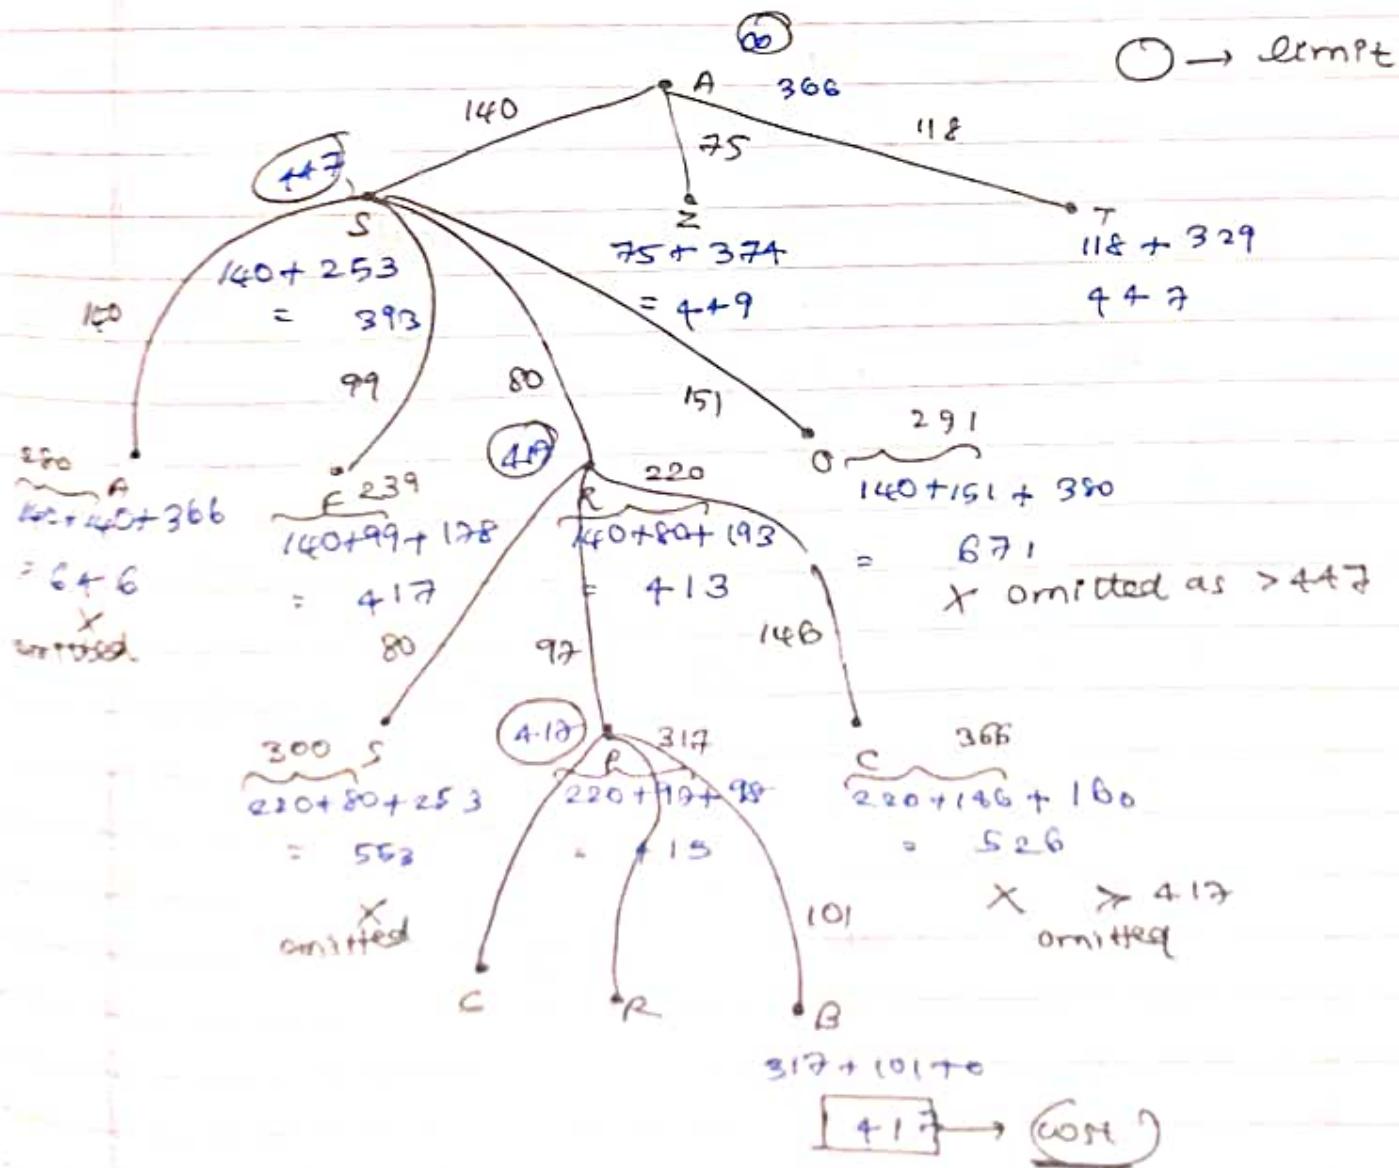
if best.f > f-limit

return failure, best.f

alternate = second lowest f-value of successor  
result, best.f = RBFS (problem, best, min(f-limit, alternate))

if result ≠ failure

return result as soln

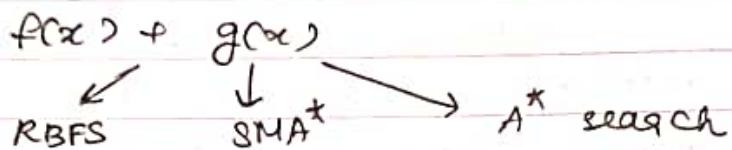


Memory bounded  $A^*$   $\Rightarrow$  MA\*

Simplified memory bounded  $A^*$   $\Rightarrow$  SMA\*



cannot add a new node without dropping an earlier node.



semi-lattice of heuristics

trivial heuristics, Dominance



Dominance :  $h_a \geq h_c$  if

$$\forall n : h_a(n) \geq h_c(n)$$

Heuristics form a semi-lattice :

max of admissible heuristics is admissible

$$ch(n) = \max [h_a(n), h_b(n)]$$

trivial heuristics

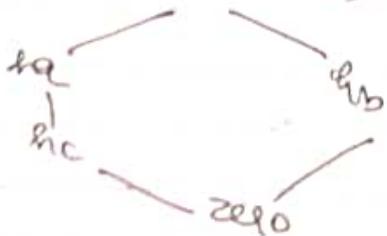
\* Bottom of lattice is the zero heuristic

(what does this give us?)

\* Top of lattice is the exact heuristic

exact

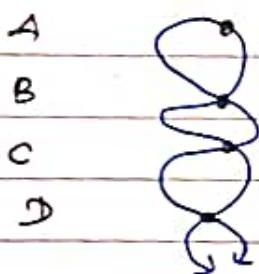
$$\max(h_a, h_b)$$



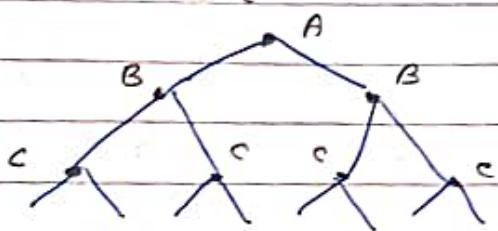
Graph search → loop.

Tree search → failure to detect repeated states can cause exponentially more work.

state graph



search tree



In BFS, don't expand the repeating node

Idea: never expand a state twice.

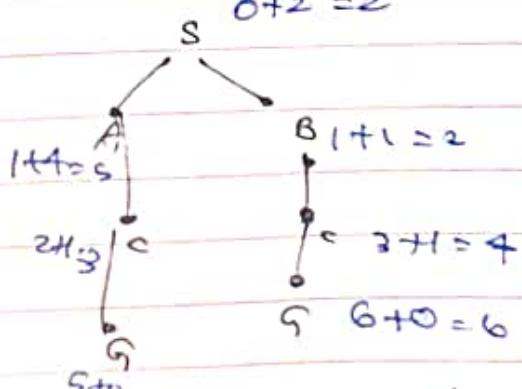
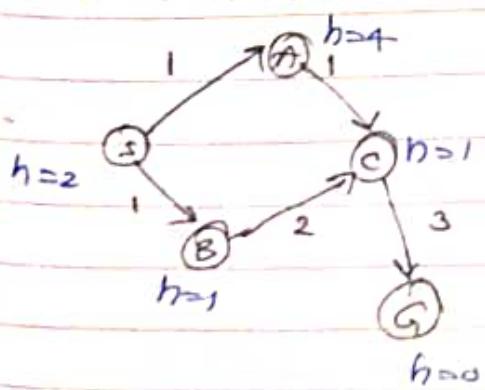
How to implement:

\* Tree search + set of expanded states  
(closed set)

- \* Expand the search tree node by node but
- \* Before expanding a node, check to make sure its state has never been expanded before
- \* If not new, skip it, if new add to closed set.

Important: store the closed set as a set,  
not as a list

A\* Graph search gone wrong?



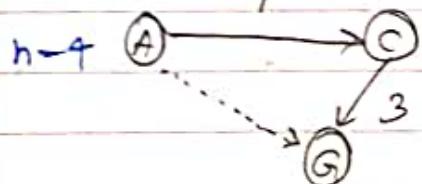
A\* → checks local maxima & minima  
but we need global maxima & minima

## consistency of heuristics

main idea: estimated heuristic costs  $\leq$  actual costs

admissibility: heuristic  $\leq$  actual goal cost

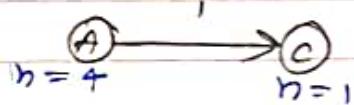
$h(A) \leq$  actual cost from A to G



consistency

heuristic cost  $\leq$  actual cost for each arc

$$h(A) - h(B) \leq \text{cost}(A \text{ to } B)$$



consequences of consistency

the f value along a path never decreases

$$h(A) \leq \text{cost}(A \text{ to } C) + h(C)$$

$A^*$  graph search is optimal

## optimality of $A^*$ graph search

sketch: consider what  $A^*$  does with a consistent heuristic

fact 1: in tree search,  $A^*$  expands nodes in increasing total f value (f-contours)

fact 2: for every state  $s$ , nodes that lead to optimally are expanded before nodes that reach  $s$  sub-optimally.

Result:  $A^*$  graph search is optimal



## optimality

### Tree search:

- \* A\* is optimal if heuristic  $h$  is admissible
- \* UDS is a special case ( $h=0$ )

### Graph search:

- \* A\* is optimal if heuristic is consistent
- \* UDS is optimal [ $h=0$  is consistent]

consistency implies admissibility.

In general, most natural admissible heuristics tend to be consistent especially if from relaxed problems.

### A\* summary

- \* A\* uses both backward costs and estimate of forward costs.
- \* A\* is optimal with admissible/inconsistent heuristics.
- \* heuristic design is key: often uses relaxed problem

### optimality of A\* Graph search



complete mathematical proof during  
next course: AI and applications

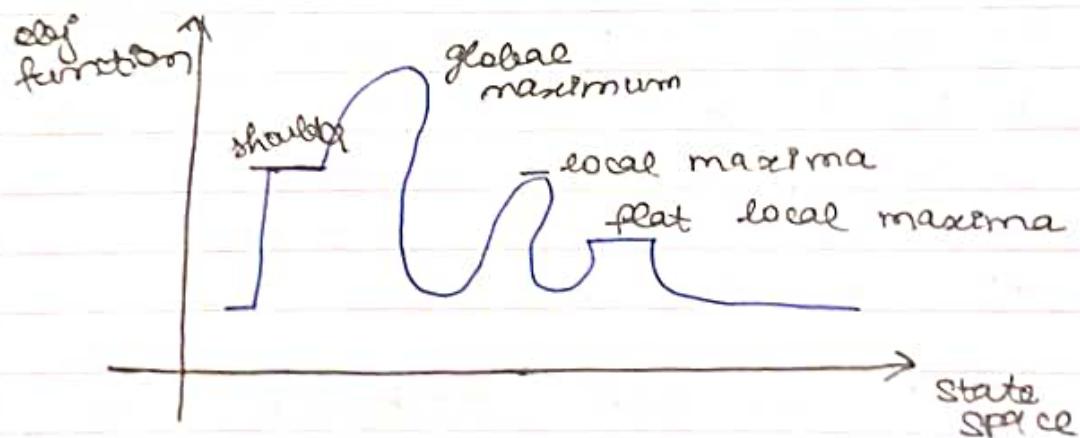
All above algorithms → Local search algorithms

### Local search algorithms

- ① Hill climbing
- ② Simulated annealing
- ③ Local beam search
- ④ Genetic algorithm search

(↓  
greedy approach)

→ current node [only local  
extremum    seems, not global]



### Hill climbing search

↓  
always keep climbing up-hill

function HC (problem)

current  $\leftarrow$  make-node (problem, initial state)

do

neighbour  $\leftarrow$  a high valued successor of current

if neighbour. value  $\leq$  current value

then current.state

current  $\leftarrow$  neighbour

Heuristic :

local maxima  $\rightarrow$  peak

Ridge : result of so many local maxima,  
available close to each other.

Plateau : flat area of state space

$\hookrightarrow$  flat local maxima / shoulder.

[no hill climbing  $\rightarrow$  try to avoid plateau as much as possible]

### 8 Queen's Problem

8 Queen move :  $2 \times 2 = 16$

taking heuristic  $f_1 = 17 \rightarrow$  1st form  
hill climbing

18	14	15	15					
14	14	13	12	12				

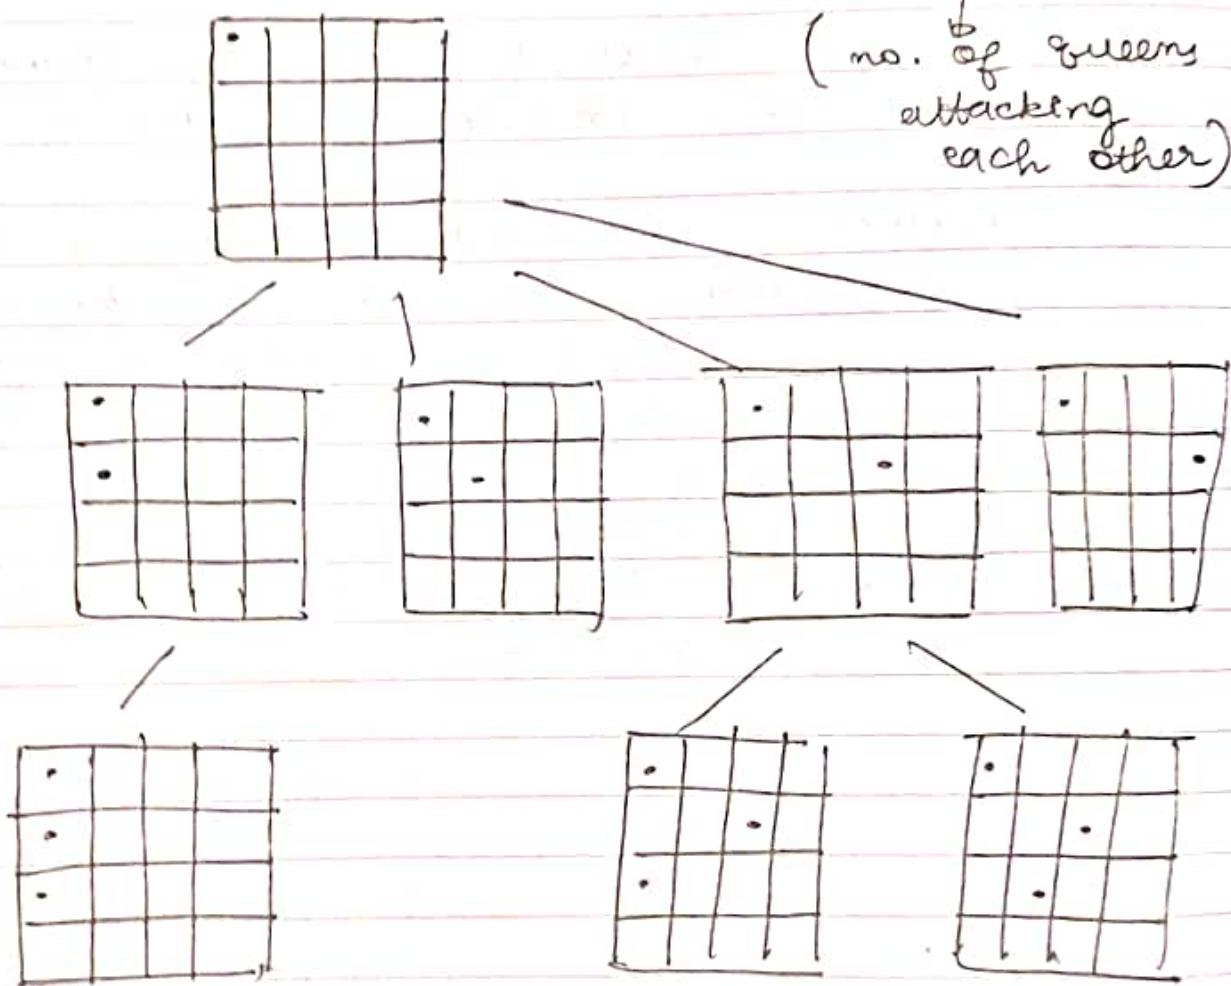
placed queen  $\rightarrow$  only go in diagonal direction  
else  $\rightarrow$  keep going sideways

move left, move right, } actions  
move diagonal, place queen } possible

(Finally place 8 Queens)

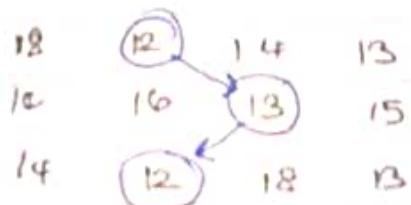
state space formulation  
4 Queen Problem

heuristic function  
(no. of queens attacking each other)



travel in the direction of least value of heuristic function

drawback: too much memory, some states can be obviously left out & not expanded



sideways } problems  
ridges }

## More search methods

- \* Local search → ④ types
- \* Local search in continuous spaces
- (decision trees) { \* searching with non-deterministic actions
- \* online search (agent is executing actions)

## Local search Algorithms and optimization problems

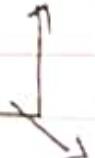


- \* complete state formulation  
for example, for 8 queens problem, all 8 Queens are on the board and need to be moved around to get to a goal state
- \* Equivalent to optimization problem often found in science and engineering.
- \* start somewhere and try to get to the solution from there
- \* local search around the current state to decide where to go next

### Pose Estimation Example

Given a geometric model of a 3D object and a 2D image of the object determine the position and orientation of the object w.r.t the camera that snapped the image.

state ( $x, y, z, x\text{-angle}, y\text{-angle}, z\text{-angle}$ )



## Hill climbing [from ppt]

Goal : optimizing an objective function  
 does not require differentiable functions.  
 can be applied to "goal" predicate type  
 of problems.

BSAT with objective function number of  
 clauses satisfied.

**Intuition:** always move to a better state

some Hill climbing Algo's

- \* start = random state or special state
- \* while (no improvement)
  - ⇒ steepest ascent: find best successor
  - ⇒ (or) greedy: select first improving successor
  - ⇒ go to that successor
- \* repeat the above process some number of times (restarts).
- \* can be done with partial solutions or full solutions

implicit in the scheme is the notion of a neighbourhood that in some way preserves the cost behaviour of the solution space.



Thinking about the TSP problem again...

If a current tour is present what would a neighbouring tour look like?

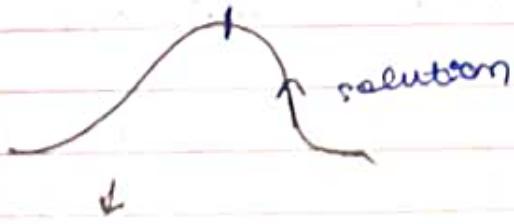


This is the way of asking for a successor function.

- \* The successor function is where the intelligence lies in hill climbing search
- \* It has to be conservative enough to preserve significant "good" portions of the current solution.
- \* and liberal enough to allow the state space to be searched without also generating into a random walk.

### Hill climbing algorithm

- \* In best first, replace storage by a single node
- \* works if single hill
- ← \* use restarts if multiple hills
- variation of hill climbing search*     \* Problems :
  - \* finds local maximum not global
  - \* plateaux: large flat regions (happens in BSAT)
  - \* ridges: fast up ridge, slow on ridge
  - \* not complete, not optimal
  - \* no memory problems



"gradient ascent"  
[ here solutions are max not min ]

Often used for numerical optimization problem

### AI Hill climbing

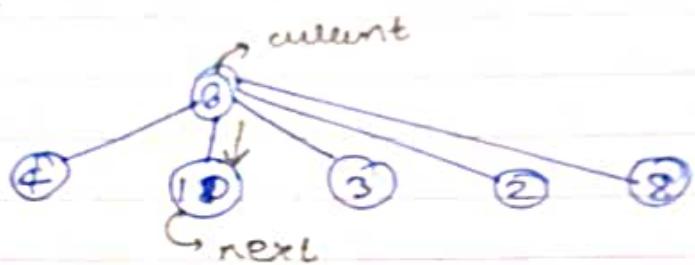
#### steepest - ascent Hill climbing

current ← start node

loop do

\* neighbor ← a highest valued successor  
of current

- \* if neighbour.value < current.value  
then return current.state
  - \* current  $\leftarrow$  neighbour
- end loop



if current = 12, again expand from  
<sup>highest</sup>  
 to (next, lowest) / not possible,  
 change heuristic function.

local maxima    ↴ .  
 Plateaus    ↴ .

Diagonal edges  
 ↴

what is it sensitive to? e.g., it  
 have any advantages?



local Maxima (minima)

- ⇒ \* Hill climbing is subject to getting stuck in a variety of <sup>local</sup> conditions,
- \* Two solutions:
  - ↳ stochastic hill climbing  $\rightarrow$  chooses random from among the uphill moves
  - $\Rightarrow$  fast stochastic hill climbing  $\rightarrow$  similar to stochastic but creates successors at random
  - $\Rightarrow$  random restart hill climbing  $\rightarrow$

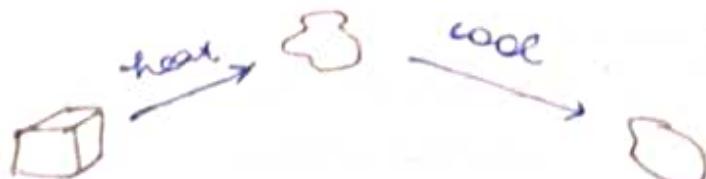
- randomly generates initial states and keep climbing high  $\rightarrow$  \* (Queens Problem)
- \* simulated annealing

### Random restart hill climbing

- \* generate a random start state
- \* run a hill climbing & store answer
- \* iterate, keeping the current best answer as you go
- \* stop when goal is reached.

### Simulated Annealing

- \* variant of hill climbing (so up & good)
- \* tries to explore enough of the search space early on, so that the final solution is less sensitive to the start state.
- \* may make some downhill moves before finding a good way to move uphill.
  - $\Rightarrow$  local maxima  $\rightarrow$  global maxima
  - $\Rightarrow$  local minima  $\rightarrow$  global minima
- \* comes from the physical process of annealing in which substances are raised to high energy levels (melted) and then cooled to solid state.



- \* the probability of moving to a higher energy state instead of lower is
 
$$P = e^{-\Delta E / kT}$$
 where  $\Delta E$  is the positive change in energy level,  $T$  is the temperature, and  $k_B$  is Boltzmann's constant.
- \* like hill climbing, but probabilistically allows down moves, controlled by current temperature & how bad move is
- \* let  $t[1], t[2] \dots$  be a temperature schedule usually  $t[1]$  is high,  $t[k] = 0.9 t[k-1]$
- \* let  $E$  be quality measure of state.  
Goal: maximize  $E$ .

### simulated Annealing algorithm

- \* current = random state,  $K=1$
- \* If  $-r(K) = 0$ , Stop
- \* next = random next state.
- \* If next is better than start, move
- \* If next is worse:
  - $\Rightarrow$  use delta,  $E(\text{next}) - E(\text{current})$
  - $\hookrightarrow$  Move to next with probability  $e^{\Delta / (K T)}$
- \*  $K = K + 1$

### simulated Annealing discussion

- \* no guarantees
- \* when  $T$  is large,  $e^{-\Delta / (K T)} \rightarrow 1$ .  
so for large  $T$ , go anywhere.

- \* when  $\tau$  is small,  $e^{\Delta E/\tau} \rightarrow e^{-\infty}$  (as) 0.  
Avoid most bad move.
- \* when  $\tau$  is zero, do simple hill climbing
- \* Execution time depends on schedule, memory use is trivial.
- \* At the beginning, the temperature is high.
- \* As the temperature becomes lower  
 $kT \downarrow \Delta E/kT \uparrow -\Delta E/kT \downarrow e^{-\Delta E/kT} \downarrow$
- \* As the process continues, the probability of a downhill move gets smaller and smaller.
- \*  $\Delta E \rightarrow$  represents the change in the value of the objective function
- \* since the physical relationships no longer apply, drop k. so  $P = e^{-\Delta E/T}$ .
- \* we need an annealing schedule, which is a sequence of values of  $\tau$ :  
 $T_0, T_1, T_2, \dots$

### Simulated Annealing Algorithm

- \* current  $\leftarrow$  start node
- \* for each  $\tau$  on the schedule
  - $\Rightarrow$  next  $\leftarrow$  randomly selected successor of current
  - $\Rightarrow$  evaluate next if it is a goal return it
  - $\Rightarrow$   $\Delta E \leftarrow$  next.value - current.value
  - $\Rightarrow$  if  $\Delta E > 0$ 
    - $\Rightarrow$  then current  $\leftarrow$  next
    - $\Rightarrow$  else current  $\leftarrow$  next with probability  $e^{-\Delta E/kT}$

## Function

function SIMULATED\_ANNEALING(problem, schedule) return a solution state

Input: problem, a problem

schedule, a mapping from time to temperature

local variables: current, a node.

next, a node

$T$ , a "temperature" controlling the probability of downward steps

current  $\leftarrow$  MAKE-NODE(INITIAL-STATE[problem])

for  $t \leftarrow 1$  to  $\infty$  do

$T \leftarrow$  schedule( $t$ )

if  $T=0$  then return current

next  $\leftarrow$  a randomly selected successor  
of current

$\Delta E \leftarrow$  VALUE[next] - VALUE[current]

if  $\Delta E > 0$  then current  $\leftarrow$  next

else current  $\leftarrow$  next with probability  $e^{\frac{-\Delta E}{T}}$

## Properties of simulated annealing search

↓

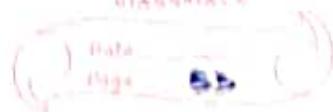
one can prove. If  $T$  decreases slowly enough, then simulated annealing search will find a global optimum with probability approaching 1.

widely used in VLSI layout, airline scheduling, etc

heuristic function

$n \rightarrow$  no. of flights on runway

- no. of flights ready to take off



Probabilistic selection

Select next with probability  $p$

o \*  $p$  = probability

random number

Generate

a random number

If its  $\leq p$  select next

Simulated Annealing procedure



At a fixed "temperature"  $T$ , state occupation probability reaches the Boltzmann distribution,

$$P(x) = \alpha e^{-E(x)/kT}$$

If  $T$  is decreased slowly enough (very slowly), the procedure will reach the best state.

Slowly enough has proven too slow for some researchers who have developed alternate schedules.

Simulated Annealing schedule



Acceptance criterion & cooling schedule



if ( $\Delta \geq 0$ ) accept

else if ( $\text{random} < e^{\Delta/T}$ ) accept

else reject

[ $0 \leq \text{random} \leq 1$ ]

(Aim: global maximum)

## simulated annealing applications

### Basic Problems

- \* Travelling salesman
- \* Graph partitioning
- \* Matching problems
- \* Graph colouring
- \* Scheduling

### Engineering

- \* VLSI design
  - ⇒ placement
  - ⇒ routing
  - ⇒ array logic minimization
  - ⇒ layout
- \* Facilities layout
- \* Image processing
- \* Code design in information theory

### Beam

- \* mix of hill climbing & best first
  - \* storage is a cache of best  $k$  states
  - \* sales steepest problem but (not optimal and not complete)
- $k$  selection,

$k$ - generated  
 @ each step → successors  
 choose best → moves in that direction



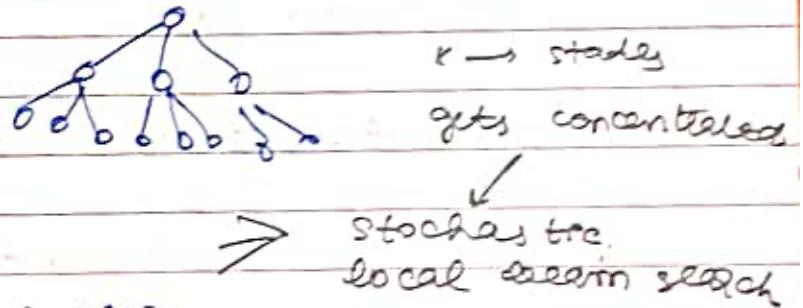
(B) any required amount

## local beam search

keeps more previous states in memory  
 simulated annealing just kept one previous state in memory.

This state keeps  $k$  states in memory.

- Randomly generates  $k$  initial states.
- If any state is a goal terminate
- Else generate all successors and select the best  $k$
- Repeat



## local (iterative) improving

Initial state = full immediate solution

Greedy hill climbing

If up, do it

If flat, probabilistically decide to accept move

If down, don't do it

We are gradually expanding the possible moves



## Local improving: performance.

- \* solves  $10^6$  queen problem quickly

- \* useful for scheduling

- \* useful for BSAT

- \* solves (sometimes) large problems

- \* more time, better answer

- \* no memory problems

- \* no guarantees of anything

## Genetic Algorithm

- \* weakly analogous to evolution
- \* no theoretic guarantees
- \* applies to nearly any problem
- \* population = set of individuals
- \* fitness function on individuals
- \* mutation operator : new individual from old one
- \* cross over : new individuals from parents

### GA Algorithm (a version)

- \* population = random set of  $n$  individuals
- \* probabilistically choose  $n$  pairs of individuals to mate
- \* probabilistically choose  $n$  descendants for next generation (may include parents or not)
- \* probability depends on fitness function & in simulated annealing
- \* How well it works  $\Rightarrow$  ?

scores to probabilities

suppose the scores of  $n$  individuals are  $a[1], a[2], \dots, a[n]$

The probability of choosing the  $j^{th}$  individual :

$$\text{prob} = \frac{a_j}{a_1 + a_2 + \dots + a_n}$$

GA Example : Boolean satisfiability

Individual → bindings for variable

Mutation → flip a variable

Cross-over → for 2 parents, randomly positions from 1 parent. For one son take those binding and use other parent for others.

Fitness → no. of clauses solved.

GA example : N-Queen's problem

eg

3-561478 individual ← array indicating column where mating ← cross over      ;<sup>th</sup> queen is assigned either (minimizes) ← no. of constraint violation

GA function optimization (Ex).

- + set  $f(x, y)$  be the function to optimize.
- + Domain for  $x$  and  $y$  is a real number between 0 and 10.
- + say the hidden function is :

$$f(x, y) = 2 \text{ if } x > 9 \& y > 9$$

$$f(x, y) = 1 \text{ if } x > 9 \text{ or } y > 9$$

$$f(x, y) = 0 \text{ otherwise}$$

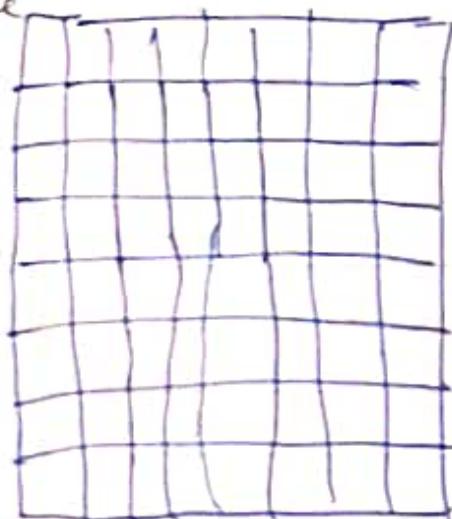
(cross over)

2	4	7	4	8	5	5	2
---	---	---	---	---	---	---	---

3	2	7	5	2	4	1	1
---	---	---	---	---	---	---	---

2	4	1	5	1	2	4
---	---	---	---	---	---	---

3	2	5	4	3	2	1	3
---	---	---	---	---	---	---	---



GA works well here.

- \* Individual = point =  $\{x, y\}$
- \* mating: something from each so:  
mate ( $\{x_1, y_1\}, \{x_2, y_2\}$ ) is  $\{x_1, y_2\} \& \{x_2, y_1\}$
- \* no mutation
- \* hill-climbing does poorly, GA does well
- \* this example generalizes function with large arity.

$$\begin{array}{ccccccccc} 3 & 2 & 7 & 5 & 2 & 4 & 1 & 1 \\ 2 & 4 & 7 & 4 & 8 & 5 & 5 & 2 \end{array} \times \begin{array}{ccccccccc} 3 & 2 & 7 & 4 & 8 & 5 & 5 & 2 \end{array}$$

### GA discussion

- \* reported to work well on some problems
- \* typically not compared with other approach  
eg hill climbing with restart.
- \* opinion: works if the "mating" operator captures good substrings.

### Genetic Algorithm

- \* start with random population of states
- \* representation serialized (string of characters or bits)
- \* states are ranked with fitness function
- \* produce new generations
  - select random pair(s) using probability
  - probability & fitness
- \* randomly choose "crossover point"
- \* offspring mix values
- \* randomly mutate

## Genetic Algorithm

- \* Given: population  $P$  and fitness function  $f$
- \* repeat
  - \* new  $P \leftarrow$  empty set
  - \* for  $i=1$  to  $\text{size}(P)$ 
    - $x \leftarrow$  random selection ( $P, \frac{1}{f}$ )
    - $y \leftarrow$  random selection ( $P, \frac{1}{f}$ )
    - child  $\leftarrow$  reproduce ( $x, y$ )
    - if (small random probability) then  
child  $\leftarrow$  mutate (child)
    - add child to new  $P$
  - \*  $P \leftarrow$  new  $P$
  - \* until some individual is fit enough or enough time has elapsed.
  - \* return the best individual in  $P$  according to  $f$ .

### Using Genetic Algorithms

- \* important aspects to use them

- ① How to encode your real life problem
- ② choice of fitness function

### Research Example

we want to generate a new operator for finding interesting points on images.

the operator will be some function of a set of values  $v_1, v_2, \dots, v_k$ .

idea: weighted sums, weighted min, weighted max,

$$v_1 \cdot a_1 + v_2 \cdot a_2 + \dots + v_k \cdot a_k = \sum v_i \cdot a_i$$

## Local search in continuous space

Given a continuous state space

$$S = \{(x_1, x_2, \dots, x_n) \mid x_i \in \mathbb{R}\}$$

Given a continuous objective function

$$f(x_1, x_2, \dots, x_n)$$

The gradient of the objective function

is a vector  $\nabla f = \frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_n}$ .

The gradient gives the magnitude and direction of the steepest slope at a point.

To find a maximum the basic idea is to set  $\nabla f = 0$

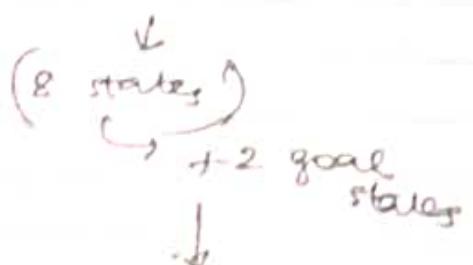
Then updating the current state become  
 $x \leftarrow x + \alpha \nabla f(x)$  where  $\alpha$  is a small constant.

Refer  $\Rightarrow$  computer vision pose

Estimation Example

searching with non-deterministic actions

vacuum world (actions = left, right, suck)



In the non-deterministic case, the result of an action can vary.

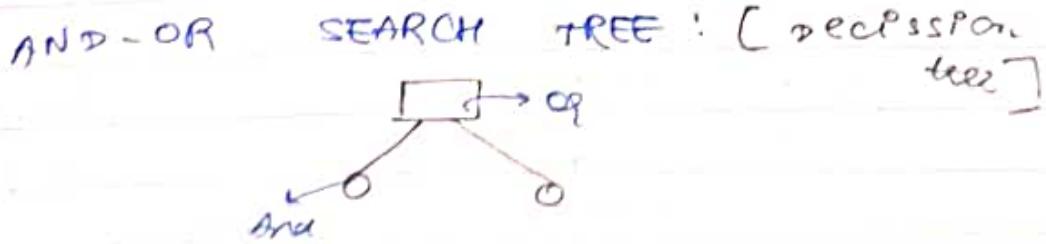
### Electric vacuum world:

when sucking a dirty square, it cleans it and sometimes, clears up dirt in an adjacent square.

when sucking a clean square, it sometimes deposits dirt on the carpet.

### Generalisation of state space model.

- ① Generalize the transition function to return a set of possible outcomes.  
 $\text{oldf} : S \times A \rightarrow S \quad \text{newf} : S \times A \rightarrow 2^S$
- ② Generalize the solution to a contingency network.  
 if state = s then action-set-1  
 else action-set-2
- ③ Generalize the search tree to an AND-OR tree.



searching with partial observations,  
 the agent does not always know its state  
 instead it maintains a belief state - a set  
 of possible states it might be in.

e.g.: a robot can be used to build a map of  
 a hostile world environment. It will have sensors  
 allowing it to see the world

Belief state space for sensorless Agent

knows if  over left

knows if on right

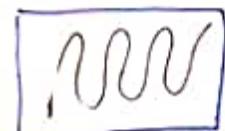
online search problems

Active agent : executes actions, acquires percents from sensors, deterministic & fully observable, has to perform an action to know the outcome.  
 Eg. web search, autonomous vehicle

Agents → core of AI



search algorithm



automatic google news

Google News :

- sports
- ML, AI
- central government
- US universities
- IIT mainly ignored

frequent  
keep up.

go down

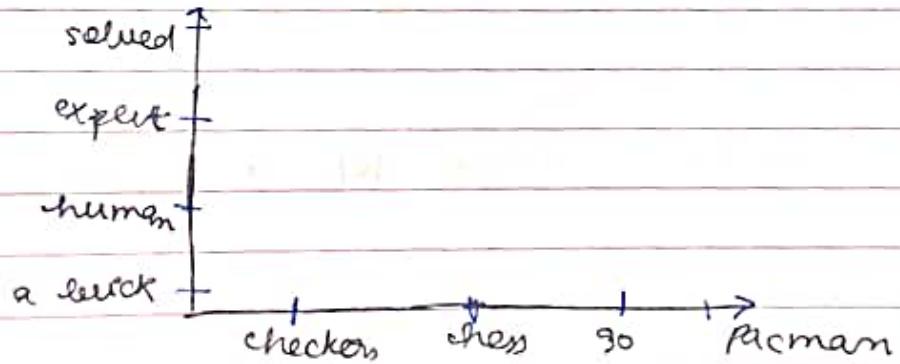
"depth limited search" +  
"local beam search" + "nearest"

best first,

← LinkedIn : employees looking for good resumes

Product Recommendations : E-commerce,  
simulated annealing / hill climbingexperience / skill set → heuristic function  
to get the good quality human resource

Game playing : state of the Art



"Adversarial Games"

Types of games

many different kinds of games!

Axes:

deterministic or stochastic?

one, two or more players

zero sum

Perfect information (is the state visible)

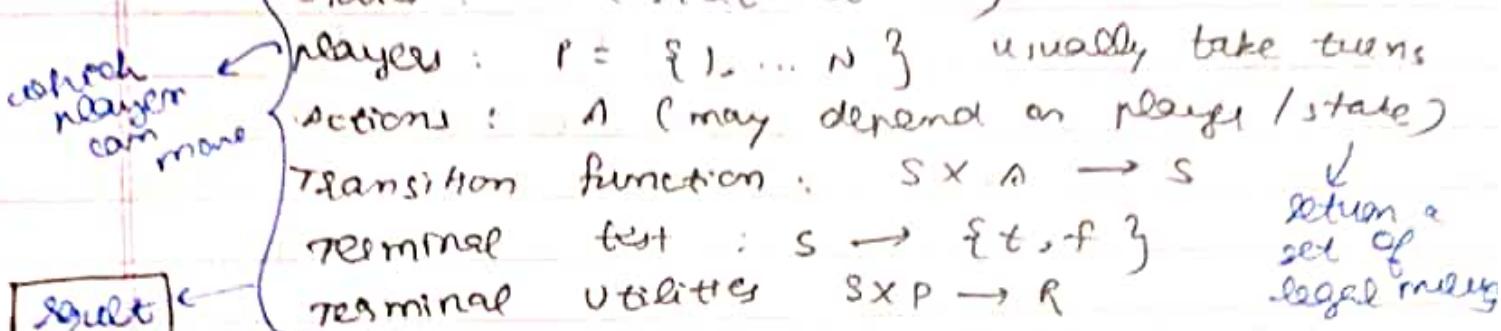
what algorithms for calculating a strategy

(policy) which recommends a move from  
each state  
↳ (search)

## Deterministic Games

Many possible formalizations, one is:

(state is start at  $s_0$ )



solution for a player is a policy  $S \rightarrow A$

true  $\rightarrow$  game is over

false  $\rightarrow$  game is not over

objective function

$\hookleftarrow$  numeric value for a game.

utility function in chess

+1 (win) -1 (loss)  $\frac{1}{2} - \frac{1}{2}$  (draw)

## Zero-sum Games

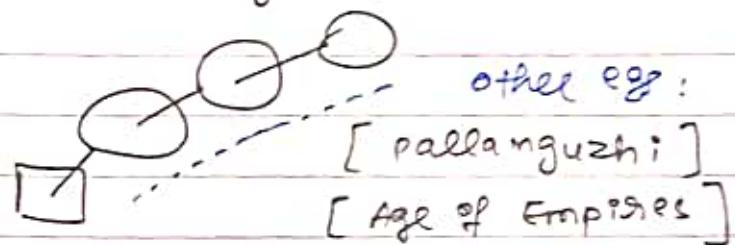
- \* Agents have opposite utilities (value on outcomes)
- \* let us think of a single value that one maximizes & the other minimizes
- \* Adversarial, pure competition  $\leftrightarrow$  chess (+1, -1, 0)

## General Games [constant sum Games]

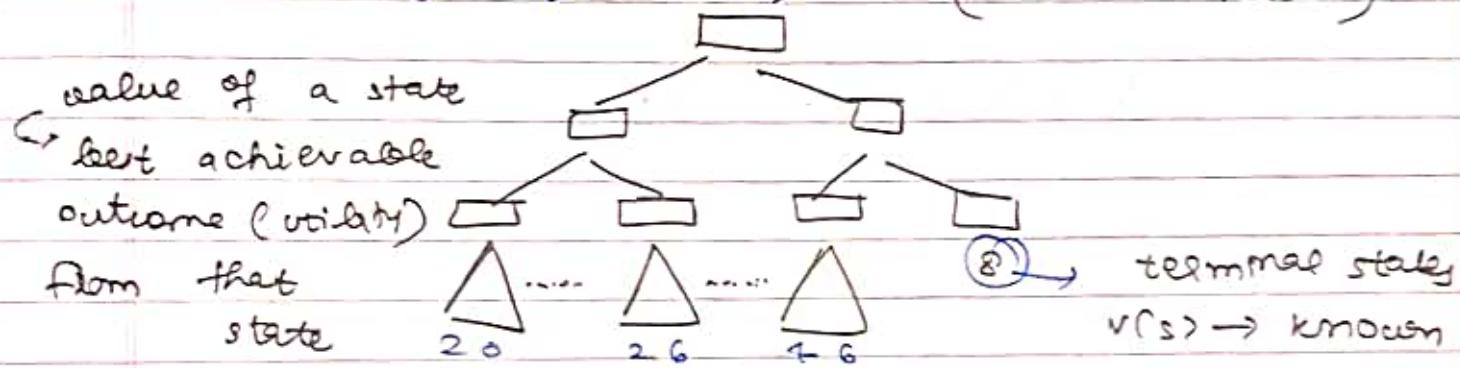
- \* Agents have independent utility (value on outcomes)
- \* cooperation, indifference, competition, and more are all possible

## Adversarial search

e.g. chess : move, thinking what other person might do



single Agent Tree → (with one reason)

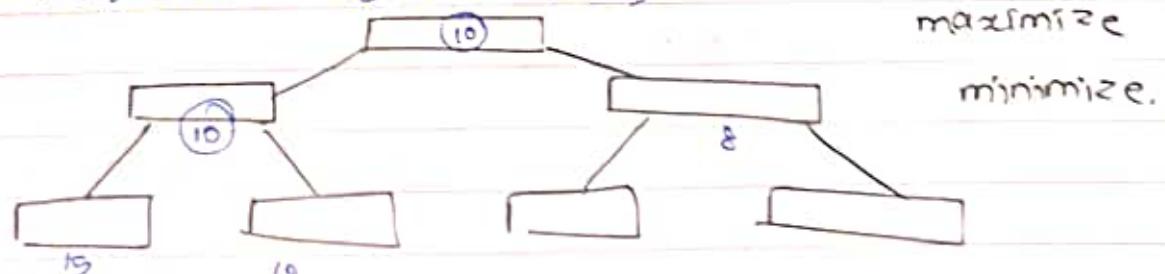


non-terminal states :

$$N(s) = \max_{s' \in \text{children}(s)} v(s')$$

(minimax algorithm)

## Adversarial Game trees



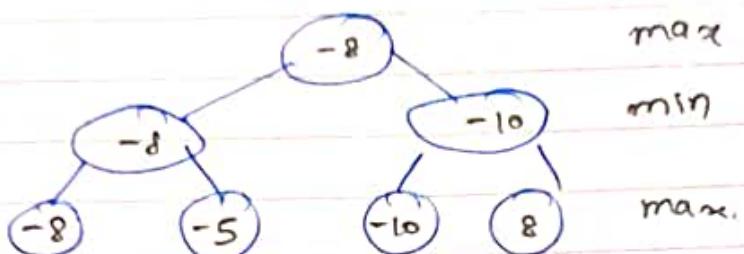
: (terminal states have a utility value)

$$N(s) = \text{known}$$

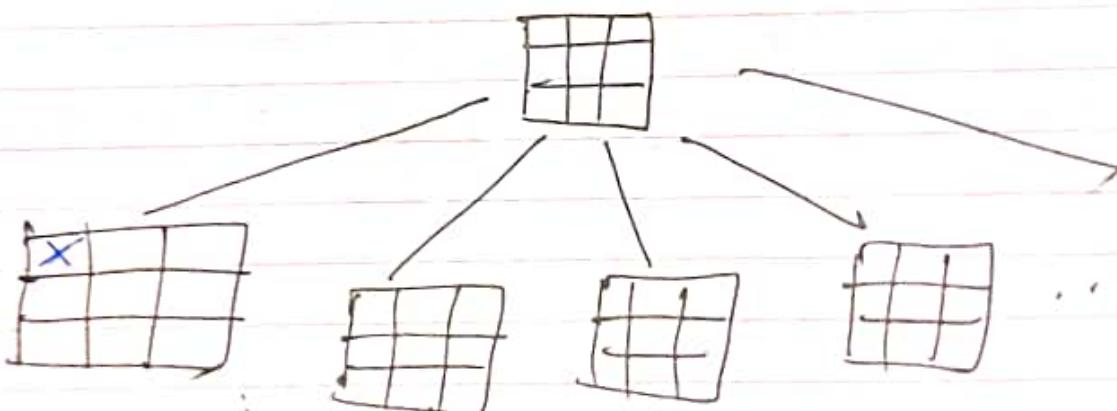
## minimax values

states under opponent's control:

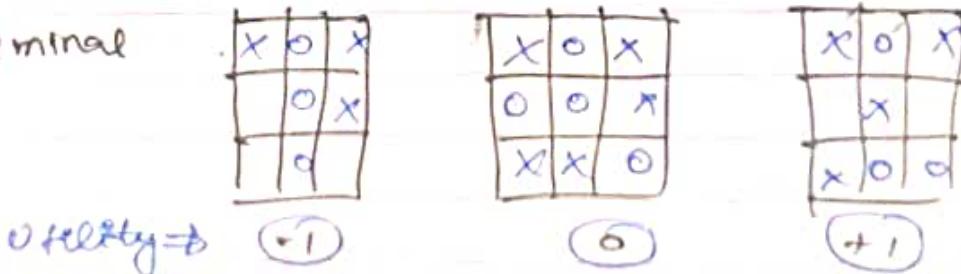
$$v(s') = \min_{\text{se successors}(s')} v(s)$$



## Tic-Tac-Toe Game



Terminal



$$9 \times 8 \times 7 \times \dots \Rightarrow 19 \text{ possible states}$$

## Adversarial search (minimax)

Deterministic, zero sum games:

- \* Tic-Tac-Toe, chess, checkers
- \* one player maximizes result
- \* the other minimizes result

## Minimax search,

- \* A state space tree.
- \* players alternate turns.
- \* compute each node's minimax value: the best achievable utility against a rational (optimal) adversary

terminal  
value,  
next  
agent

8 2 5 6

## Minimax Implementation

```
def max-value (state):
```

    Initialize  $v = -\infty$

$v(s) = \max_{s' \text{successor}} v(s')$  for each successor of state:

$v = \max(v, \min-value(\text{successors}))$

return  $v$

```
def min-value (state):
```

    Initialize  $v = +\infty$

$v(s) = \min_{s' \text{successor}} v(s')$  for each successor of state:

$v = \min(v, \max-value(\text{successors}))$

return  $v$

↓  
dispatch

```
def value (state):
```

    if the state is a terminal state:

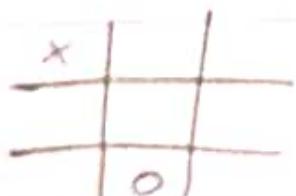
        return the state's utility

    if the next agent is MAX:

        return max-value (state)

    if the next agent is MIN:

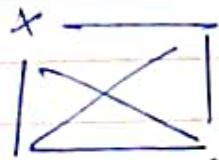
        return min-value (state)



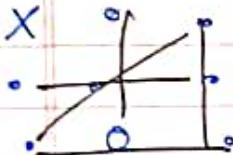
(eg)



considering

$$\begin{matrix} X & \cdot & \cdot \\ \cdot & \cdot & \cdot \\ \cdot & 0 & \cdot \end{matrix}$$


6 possibilities of success of X



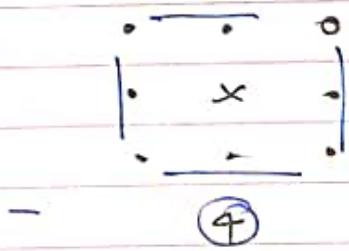
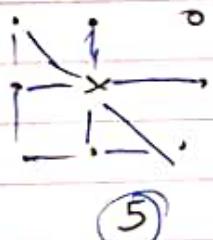
5 possibilities of success of O

$$\text{value of state} = 6 - 5 = 1$$

considering

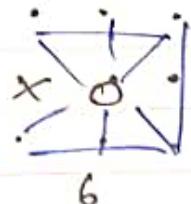
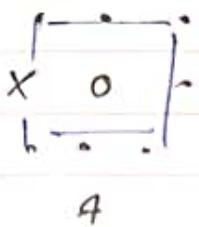
$$\begin{matrix} \cdot & \cdot & 0 \\ \cdot & X & \cdot \\ \cdot & \cdot & \cdot \end{matrix}$$

value of each state in tic tac toe can hence be calculated



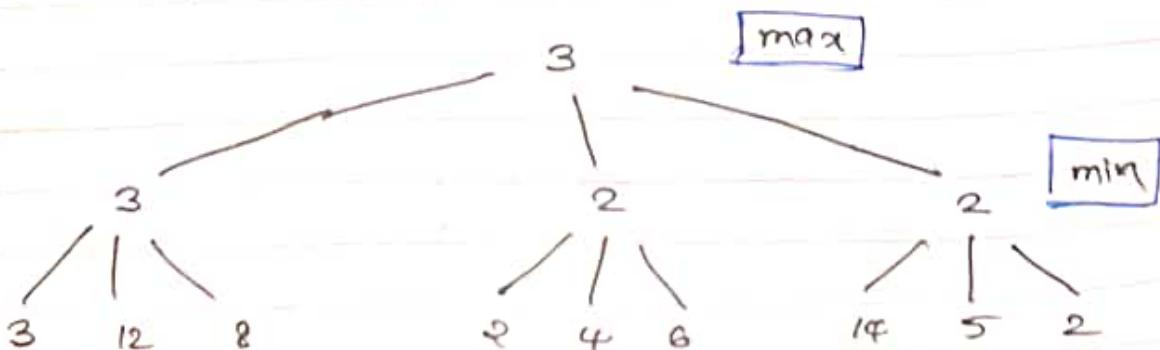
$$1 - 1 = 0$$

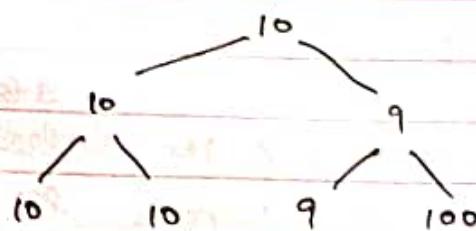
considering

$$\begin{matrix} \cdot & \cdot & \cdot \\ X & 0 & \cdot \\ \cdot & \cdot & \cdot \end{matrix}$$


$$4 - 6 = -2$$

minmax example





optimal against a perfect player.  
otherwise?

### Minmax efficiency

- \* just like exhaustive DFS
- \* time :  $O(b^m)$
- \* space :  $O(bm)$

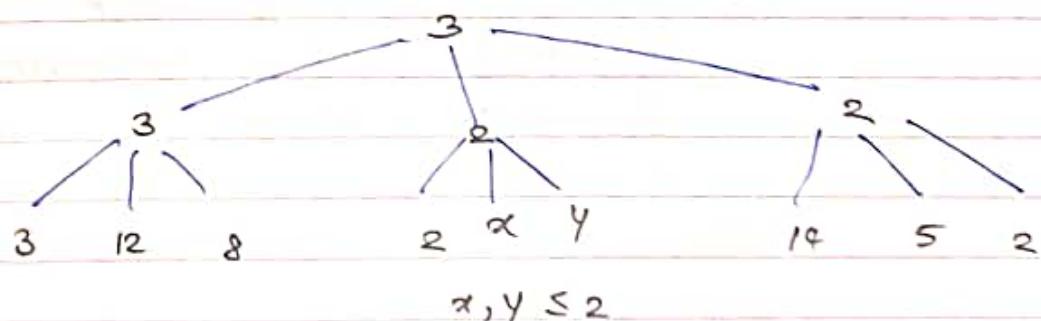
Eg : chess,  $b \approx 35$ ,  $m \approx 100$

exact solution is completely infeasible

But do we need to explore the whole tree?

### GAME TREE PRUNING

#### Minmax Pruning



#### Alpha-Beta Pruning

General configuration (MIN version) :

- \* we are computing the MIN-VALUE at some node  $n$ .
- \* we're looping over  $n$ 's children.
- \*  $n$ 's estimate of the children's min is dropping
- \* who cares about  $n$ 's value? MAX
- \* let  $a$  be the best value that MAX can get any choice point along the current path from the root.

- \* If  $n$  becomes worse than  $a$ , MAX will avoid it, so we can stop considering  $n$ 's other children (it's already bad enough that it won't be played)

[MAX version is symmetric]

alpha → minimizing

beta → maximizing

### Alpha-Beta Implementation

$\alpha$ : MAX's best option on path to root

$\beta$ : MIN's best option on path to root

```
def max_value(state,  $\alpha$ ,  $\beta$ ):
```

initialize  $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

if  $v \geq \beta$  return  $v$

$\alpha = \max(\alpha, v)$

return  $v$

```
def min_value(state,  $\alpha$ ,  $\beta$ ):
```

initialize  $v = +\infty$

for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$

if  $v \leq \alpha$  return  $v$

$\beta = \min(\beta, v)$

return  $v$

### Alpha-Beta Pruning Prevette

- \* This pruning has no effect on minimax value computed for the root.

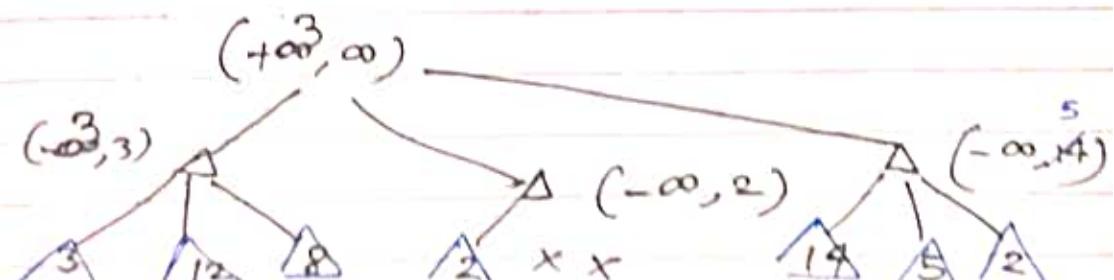
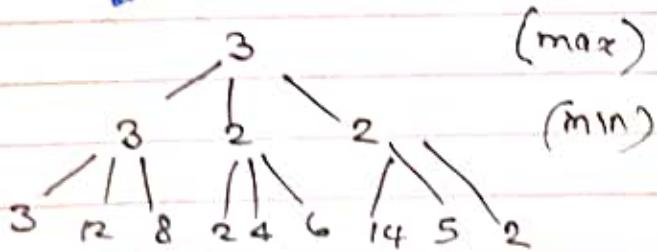
\* values of intermediate nodes might be wrong  
 ↓

Important: children of the root may have the wrong value.  
 so the most naive version won't let you do action selection

\* Good child ordering improves effectiveness of pruning.

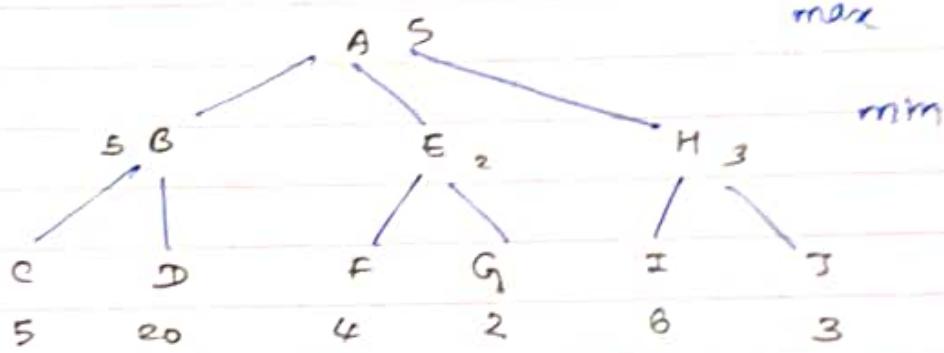
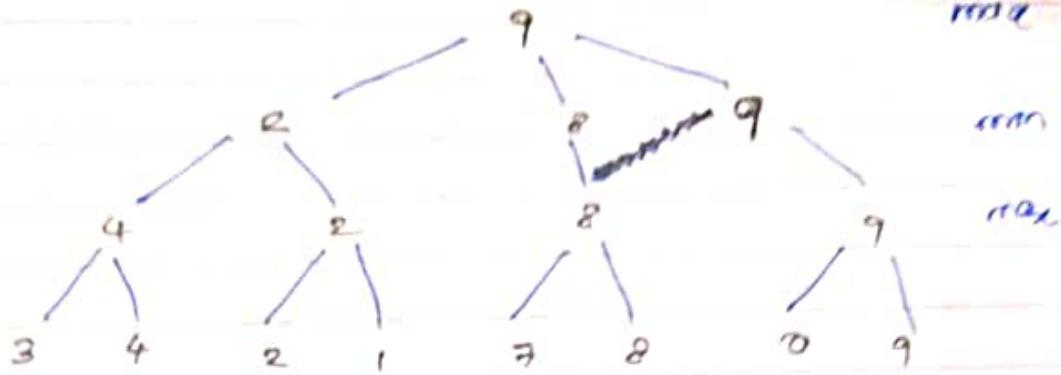
\* with "perfect ordering"  
 time complexity drops to  $O(b^{m/2})$   
 Doubts, solvable depth  
 full search of eg chess is still hopeless

α-β Pruning example

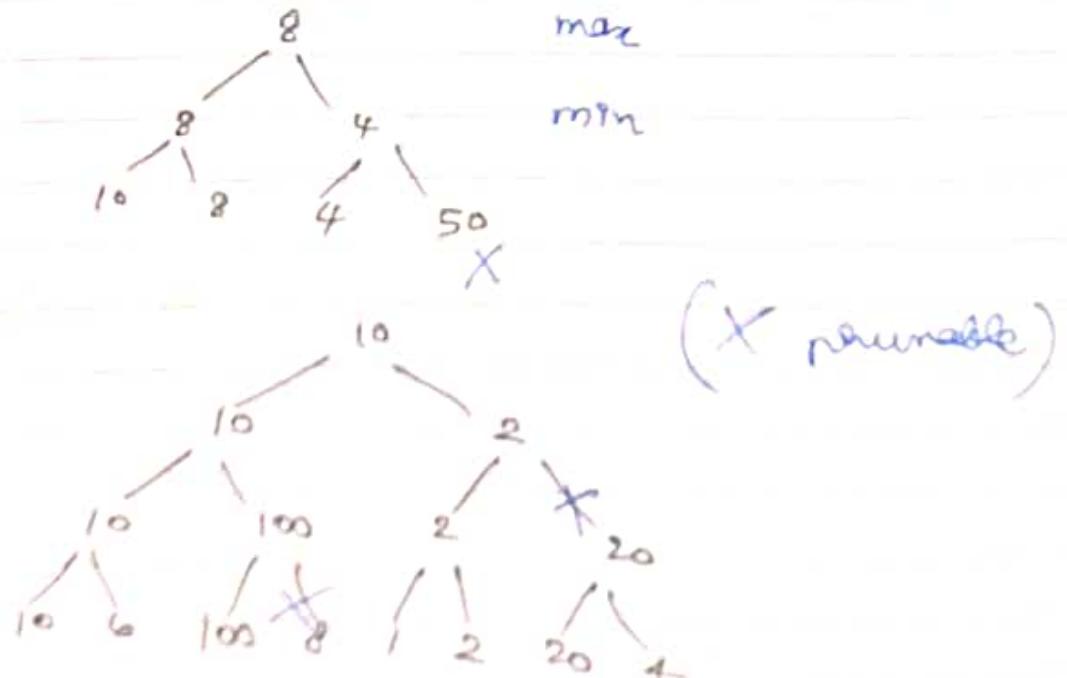


[ 2 nodes are pruned  
 by observing 9<sub>b</sub> neighbours ]

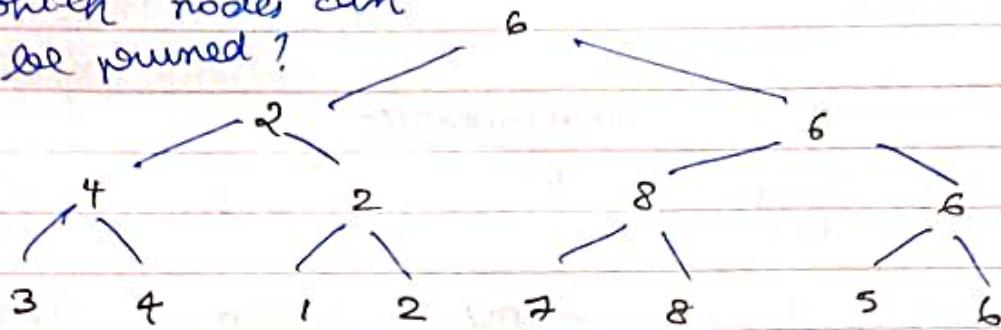
function α-β(state) returns an action  
 v → max\_value(state, -∞, ∞)  
 . return the actions in actions(state).value v



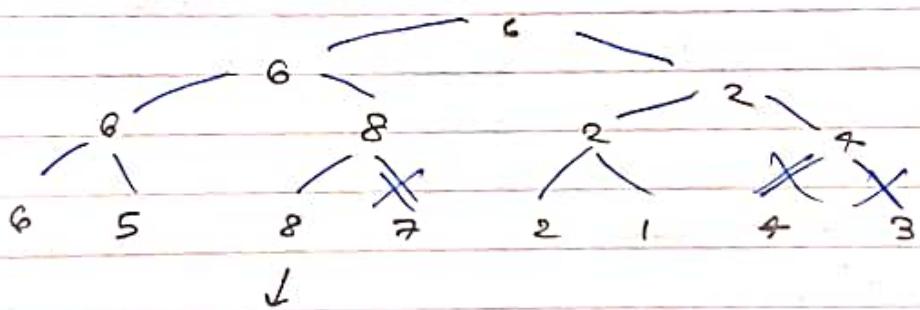
[ Perfect decision games  
 chess, checkers, monopoly,  
 Imperfect decision games  
 card games ]



which nodes can be pruned?



↓  
no node can be (expanded) pruned  
as the most favourable nodes for both  
are explored last (i.e. in the diagram,  
are on the right-hand side)



↓  
lots of nodes can be pruned as the  
most favourable nodes for both are  
explored first (i.e. in the diagram are  
on the left-hand side).

killer move → Best moves are called as killer moves.

transposition ↓ → repeated states occurring frequently

transposition } moves → moving! eg checkers,  
table } panogale no solitaire

( TIC-TAC  
- TOE also )

## Games

	deterministic	stochastic
perfect information	chess, checkers, go	backgammon, monopoly
imperfect information	blond tic-tac-toe, battleship	budget poker, scrabble, nullop, way

## Resource limits



- \* In real-life games, cannot search to leaves.
- \* SLM: Depth limited search  
Instead, only search to a limited depth  
in the tree  
Replace terminal utility with an evaluation  
function in non-terminal positions.

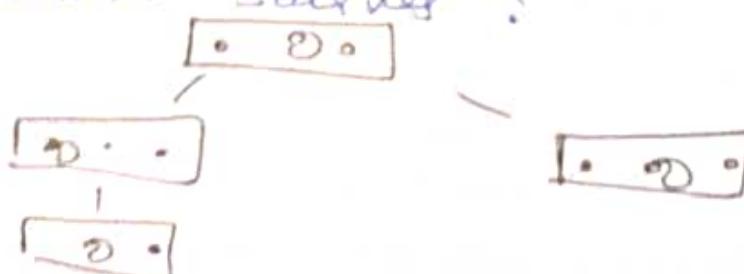
← eg

suppose we have 100 seconds, we can  
explore 10k nodes / sec.  
so can check 1M nodes per move  
 $\alpha\text{-}\beta$  reaches about depth 8 - decent  
chess program



- \* Guarantee of optimal play is gone
- \* Use iterative deepening for an  
anytime algorithm.

Why PGM stays?



A danger of replanning agent!

- \* He knows his score will go up by eating the dot now (W, E)
- \* He knows his score will go up just as much by eating the dot later (WE)
- \* There are no point scoring opportunities after eating the dot (within the horizon, two tree)
- \* Therefore, waiting seems just as good as eating: he may go east, then back west in the next round of replanning

### Evaluation Functions

↳ imperfect real time decisions

- \* minimax algorithm generates the entire game tree
- \*  $\alpha\text{-}\beta$  pruning still has to search all the way to the terminal states
- \* algorithms should cut-off the search earlier and apply heuristic evaluation function. Turns non-terminal nodes to terminal leaves
- \* replace utility function by a heuristic function which estimates the position utility and replace the terminal test by a cut off test
- \* returning an estimate of the expected utility of the game from any given position
- \* order the terminal states in the same way as the utility function
- \* computation must not take too long

cut-off test / terminal test

$\alpha$ - $\beta$  search

$v \rightarrow \text{max-value}(\text{state}, -\infty, \infty)$

$\left\{ \begin{array}{l} \text{max-value}(\text{state}, \alpha, \beta) \\ \text{if terminal-test}(\text{state}) \\ \quad \text{return utility}(\text{state}) \\ v = -\infty \\ \max(v, \text{min-value}(\text{successor})) \end{array} \right.$

$\left\{ \begin{array}{l} H\text{-minimax}(s, d) : \\ \text{EVAL}(s) \quad \text{if cut-off test}(s, d) \\ \max \quad H\text{-minimax}(\text{result}(s, a), d+1) \quad \begin{array}{l} \text{if player} \\ \text{= max} \end{array} \\ \min \\ \text{if player} \\ \text{= min} \end{array} \right.$

evaluation function  $\rightarrow$  features

TIC TAC TOE

no. of  $X$  is your feature,

no. of  $X$  is a column row

$X \ X \ ?$

$X$   
 $X$

weightage

Evaluation Functions

Evaluation functions score non-terminals in depth limited search

Ideal function: returns the actual minimax value of the position

In practice: typically weighted linear sum of features

P. Eval(s) =  $w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$   
no threshold

## depth matter

- \* Evaluation functions are always imperfect.
- \* The deeper in the tree the evaluation function is buried, the less the quality of the evaluation function matters.
- \* An important example of the trade-off between complexity of features and complexity of computation.

## { Quiescent search :

↳ unlikely to exhibit wild swings.  
combination of moves

eval fn : cut off test ( $S_j$ )

## Horizon effect : ↳ (try deferring)

↳ very difficult

(eval function) program is facing opponent's move that causes a lot of damage.

Avoid horizon effect : singular extension : move clearly better than all other moves

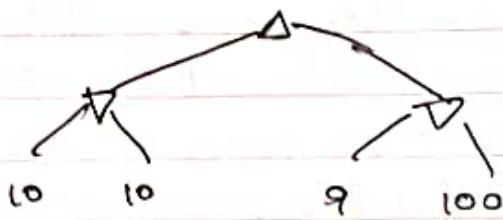
## Forward pruning

(beam search) ↳ prune a move at a given node immediately without even considering it.

↳ k-paths → best among the 'k' paths  
probabilistic cut → forward pruning using previous experience.

Uncertain outcomes

worst case vs average case



Idea: uncertain outcome, controlled by chance not an adversary?

Expectimax search



why wouldn't we know what the result of an action will be?

- \* Explicit randomness: rolling dice
- \* Unpredictable opponents: the ghosts respond randomly
- \* Actions can fail
  - ↳ while moving a robot, wheels might slip



values should now reflect average-case (expectimax) outcome, not worst-case (minimax) outcome.

Expectimax Pseudocode

def exp-value ( state ) :

    initialize  $v = \infty$

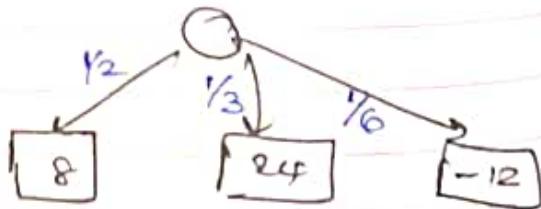
    for each successor of state:

$P = \text{probability}(\text{successor})$

$v += P * \text{value}(\text{successor})$

    return  $v$

$$v = (y_2)(8) + (y_3)(24) \\ + \frac{1}{6}(-12) = 10$$



Expectimax search : compute the average score under optimal play

- \* max nodes as in minimax search
- \* chance nodes are like min nodes but the outcome is uncertain.
- \* calculate their expected utility, i.e. take weighted average (expectation) of children

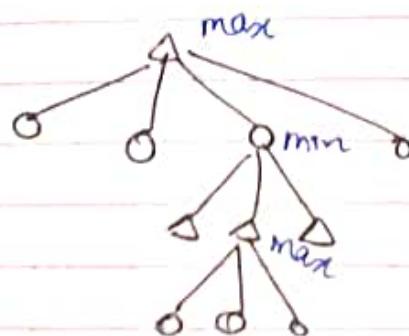
### Better Pseudocode

```

def max-value (state):
    initialize v = -∞
    for each successor of state:
        v = max (v, value (successor))
    return v

def value (state):
    if the state is a terminal state:
        return the state's utility
    if the next agent is MAX:
        return max-value (state)
    if the next agent is EXP:
        return exp-value (state)

```

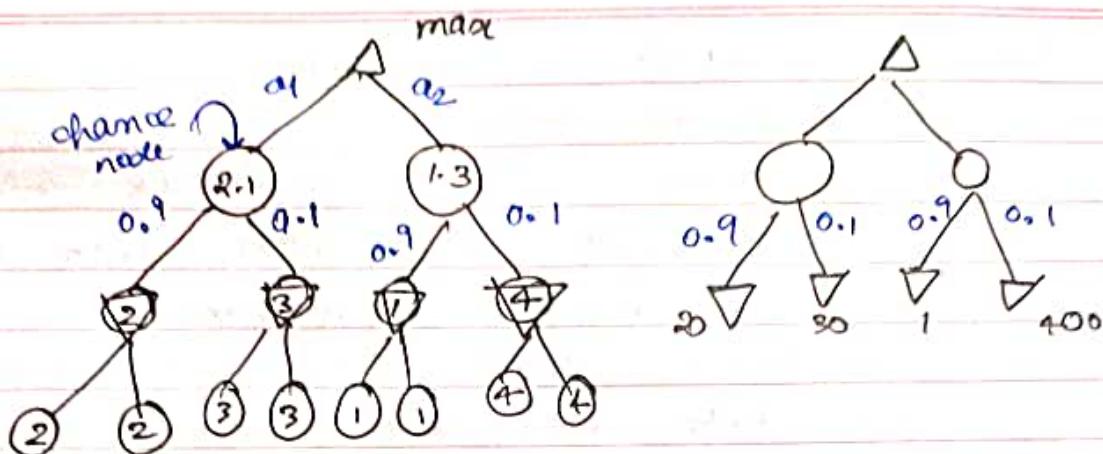


EXPECTMINIMAX (S)

UTILITY (s) IF TERMINAL TEST (s)

max ( EXPECTMINIMAX  
(Result (s, a)) IF NOT a  
MAX  
min ( ) ...

Exp (r) EX... Result (s, a)  
IF NOT chance



### Depth-Limited Expectimax

↪ @ depth

estimate of true expectimax value

(which would require a lot of work to compute)

Reminder : Probabilities -



A random variable represents an event whose outcome is unknown.

A probability distribution is an assignment of weights to outcomes.



e.g.: traffic on freeway

Random variable, T - whether there's traffic  
outcome: T in {none, light, heavy}

Distribution:  $P(T = \text{none}) = 0.25$

$P(T = \text{light}) = 0.50$ ,  $P(T = \text{heavy}) = 0.25$

Some laws of probability:

- \* probabilities are always non-negative
- \* probabilities over all possible outcomes sum to one

## other Game types

### Mixed layer types

e.g. Backgammon

### expectiminimax

- \* environment is an extra "random agent" player that moves after each min/max agent
- \* Each node computes the appropriate combination of 9 & children

### Example : Backgammon

Dice rolls increase b : 21 possible rolls with 2 dice

Backgammon  $\approx$  20 legal moves

$$\text{depth } 2 = 20 \times (21 \times 20)^3 = 1.2 \times 10^9$$



As depth increases, probability of reaching a given search node shrinks

- \* so usefulness of search is diminished
- \* so limiting depth is less damaging
- \* But pruning is trickier

### Historic AI

→ TDGammmon uses depth-2 search + very good evaluation function + reinforcement learning : world-level champion play



, so AI world champion  
in any game!

## multi-agent utilities

when game is not zero-sum or has multiple players?

Generalization of minimax:

- \* terminals have utility tuples
- \* node values are also utility tuples
- \* Each player maximizes its own component
- \* can give rise to cooperation and competition dynamically